



Neural Network Design & Deployment

Olivier Bichler, David Briand, Victor Gacoin,
Benjamin Bertelone, Thibault Allenet, Johannes C. Thiele

Friday 6th September, 2019

Contents

1	Presentation	6
1.1	Database handling	6
1.2	Data pre-processing	6
1.3	Deep network building	7
1.4	Performances evaluation	8
1.5	Hardware exports	8
1.6	Summary	10
2	About N2D2-IP	11
3	Performing simulations	11
3.1	Obtaining the latest version of this manual	11
3.2	Minimum system requirements	11
3.3	Obtaining N2D2	12
3.3.1	Prerequisites	12
	Red Hat Enterprise Linux (RHEL) 6	12
	Ubuntu	12
	Windows	13
3.3.2	Getting the sources	13
3.3.3	Compilation	13
3.4	Downloading training datasets	14
3.5	Run the learning	14
3.6	Test a learned network	14
3.6.1	Interpreting the results	14
	Recognition rate	14
	Confusion matrix	14
	Memory and computation requirements	14
	Kernels and weights distribution	14
	Output maps activity	15
3.7	Export a learned network	15
3.7.1	C export	17
3.7.2	CPP_OpenCL export	18
3.7.3	CPP_TensorRT export	19
3.7.4	CPP_cuDNN export	20
3.7.5	C_HLS export	21
3.7.6	Layer compatibility table	21
4	INI file interface	22
4.1	Syntax	22
4.1.1	Properties	22
4.1.2	Sections	22
4.1.3	Case sensitivity	22
4.1.4	Comments	22
4.1.5	Quoted values	22
4.1.6	Whitespace	22
4.1.7	Escape characters	22
4.2	Template inclusion syntax	23
4.2.1	Variable substitution	23
4.2.2	Control statements	23
	block	24

	for	24
	if	24
	include	24
4.3	Global parameters	24
4.4	Databases	24
4.4.1	MNIST	24
4.4.2	GTSRB	24
4.4.3	Directory	25
	<i>Speech Commands Dataset</i>	27
4.4.4	Other built-in databases	27
	Actitracker_Database	27
	CIFAR10_Database	27
	CIFAR100_Database	27
	CKP_Database	27
	Caltech101_DIR_Database	28
	Caltech256_DIR_Database	28
	CaltechPedestrian_Database	28
	Cityscapes_Database	29
	Daimler_Database	29
	DOTA_Database	29
	FDDB_Database	29
	GTSDB_DIR_Database	29
	ILSVRC2012_Database	30
	KITTI_Database	30
	KITTI_Road_Database	30
	KITTI_Object_Database	30
	LITISRouen_Database	30
4.4.5	Dataset images slicing	31
4.5	Stimuli data analysis	31
4.5.1	Zero-mean and unity standard deviation normalization	31
4.5.2	Subtracting the mean image of the set	32
4.6	Environment	33
4.6.1	Built-in transformations	34
	AffineTransformation	35
	ApodizationTransformation	36
	ChannelExtractionTransformation	36
	ColorSpaceTransformation	36
	DFTTransformation	37
	DistortionTransformation	37
	EqualizeTransformation	37
	ExpandLabelTransformation	38
	FilterTransformation	38
	FlipTransformation	39
	GradientFilterTransformation	39
	LabelSliceExtractionTransformation	39
	MagnitudePhaseTransformation	40
	MorphologicalReconstructionTransformation	41
	MorphologyTransformation	41
	NormalizeTransformation	42
	PadCropTransformation	42
	RandomAffineTransformation	43
	RangeAffineTransformation	45

N2D2 IP only

N2D2 IP only

N2D2 IP only

	RangeClippingTransformation	45
	RescaleTransformation	45
	ReshapeTransformation	45
	SliceExtractionTransformation	45
	ThresholdTransformation	46
	TrimTransformation	46
	WallisFilterTransformation	47
4.7	Network layers	47
4.7.1	Layer definition	47
4.7.2	Weight fillers	47
	ConstantFiller	47
	HeFiller	48
	NormalFiller	48
	UniformFiller	48
	XavierFiller	49
4.7.3	Weight solvers	49
	SGDSolver_Frame	49
	SGDSolver_Frame_CUDA	50
	AdamSolver_Frame	50
	AdamSolver_Frame_CUDA	50
4.7.4	Activation functions	50
	Logistic	51
	LogisticWithLoss	51
	Rectifier	51
	Saturation	51
	Softplus	51
	Tanh	51
	TanhLeCun	51
4.7.5	Anchor	51
	Configuration parameters (<i>Frame</i> models)	52
	Outputs remapping	53
4.7.6	Conv	53
	Configuration parameters (<i>Frame</i> models)	55
	Configuration parameters (<i>Spike</i> models)	56
4.7.7	Deconv	56
	Configuration parameters (<i>Frame</i> models)	58
4.7.8	Pool	58
	Maxout example	59
	Configuration parameters (<i>Spike</i> models)	60
4.7.9	Unpool	61
4.7.10	ElemWise	62
	Sum operation	62
	AbsSum operation	62
	EuclideanSum operation	62
	Prod operation	62
	Max operation	62
	Examples	63
4.7.11	FMP	63
	Configuration parameters (<i>Frame</i> models)	63
4.7.12	Fc	63
	Configuration parameters (<i>Frame</i> models)	64
	Configuration parameters (<i>Spike</i> models)	64

4.7.13	Rbf	65
	Configuration parameters (<i>Frame</i> models)	65
4.7.14	Softmax	66
4.7.15	LRN	66
	Configuration parameters (<i>Frame</i> models)	67
4.7.16	LSTM	67
	Global layer parameters (<i>Frame_CUDA</i> models)	67
	Configuration parameters (<i>Frame_CUDA</i> models)	68
	Current restrictions	68
	Further development requirements	68
	Development guidance	69
4.7.17	Dropout	70
	Configuration parameters (<i>Frame</i> models)	70
4.7.18	Padding	70
4.7.19	Resize	70
	Configuration parameters	71
4.7.20	BatchNorm	71
	Configuration parameters (<i>Frame</i> models)	71
4.7.21	Transformation	71
5	Tutorials	73
5.1	Learning deep neural networks: tips and tricks	73
5.1.1	Choose the learning solver	73
5.1.2	Choose the learning hyper-parameters	73
5.1.3	Convergence and normalization	74
5.2	Building a classifier neural network	74
5.3	Building a segmentation neural network	77
5.3.1	Faces detection	79
5.3.2	Gender recognition	79
5.3.3	ROIs extraction	80
5.3.4	Data visualization	81
5.4	Transcoding a learned network in spike-coding	81
5.4.1	Render the network compatible with spike simulations	81
5.4.2	Configure spike-coding parameters	82

1 Presentation

The N2D2 platform is a comprehensive solution for fast and accurate Deep Neural Network (DNN) simulation and full and automated DNN-based applications building. The platform integrates database construction, data pre-processing, network building, benchmarking and hardware export to various targets. It is particularly useful for DNN design and exploration, allowing simple and fast prototyping of DNN with different topologies. It is possible to define and learn multiple network topology variations and compare the performances (in terms of recognition rate and computational cost) automatically. Export targets include CPU, DSP and GPU with OpenMP, OpenCL, Cuda, cuDNN and TensorRT programming models as well as custom hardware IP code generation with High-Level Synthesis for FPGA and dedicated configurable DNN accelerator IP¹.

In the following, the first section describes the database handling capabilities of the tool, which can automatically generate learning, validation and testing data sets from any hand made database (for example from simple files directories). The second section briefly describes the data pre-processing capabilities built-in the tool, which does not require any external pre-processing step and can handle many data transformation, normalization and augmentation (for example using elastic distortion to improve the learning). The third section show an example of DNN building using a simple INI text configuration file. The fourth section show some examples of metrics obtained after the learning and testing to evaluate the performances of the learned DNN. Next, the fifth section introduces the DNN hardware export capabilities of the toolflow, which can automatically generate ready to use code for various targets such as embedded GPUs or full custom dedicated FPGA IP. Finally, we conclude by summarising the main features of the tool.

1.1 Database handling

The tool integrates everything needed to handle custom or hand made databases:

- Genericity: load image and sound, 1D, 2D or 3D data;
- Associate a label for each data point (useful for scene labeling for example) or a single label to each data file (one object/class per image for example), 1D or 2D labels;
- Advanced Region of Interest (ROI) handling:
 - Support arbitrary ROI shapes (circular, rectangular, polygonal or pixelwise defined);
 - Convert ROIs to data point (pixelwise) labels;
 - Extract one or multiple ROIs from an initial dataset to create as many corresponding additional data to feed the DNN;
- Native support of file directory-based databases, where each sub-directory represents a different label. Most used image file formats are supported (JPEG, PNG, PGM...);
- Possibility to add custom datafile format in the tool without any change in the code base;
- Automatic random partitionning of the database into learning, validation and testing sets.

1.2 Data pre-processing

Data pre-processing, such as image rescaling, normalization, filtering... is directly integrated into the toolflow, with no need for external tool or pre-processing. Each pre-processing step is called a *transformation*.

The full sequence of transformations can be specified easily in a INI text configuration file. For example:

```
; First step: convert the image to grayscale
[env.Transformation-1]
Type=ChannelExtractionTransformation
CSChannel=Gray
```

¹Ongoing work

```

; Second step: rescale the image to a 29x29 size
[env.Transformation-2]
Type=RescaleTransformation
Width=29
Height=29

; Third step: apply histogram equalization to the image
[env.Transformation-3]
Type=EqualizeTransformation

; Fourth step (only during learning): apply random elastic distortions to the images to extent the
learning set
[env.OnTheFlyTransformation]
Type=DistortionTransformation
ApplyTo=LearnOnly
ElasticGaussianSize=21
ElasticSigma=6.0
ElasticScaling=20.0
Scaling=15.0
Rotation=15.0

```

Example of pre-processing transformations built-in in the tool are:

- Image color space change and color channel extraction;
- Elastic distortion;
- Histogram equalization (including CLAHE);
- Convolutional filtering of the image with custom or pre-defined kernels (Gaussian, Gabor...);
- (Random) image flipping;
- (Random) extraction of fixed-size slices in a given label (for multi-label images)
- Normalization;
- Rescaling, padding/cropping, trimming;
- Image data range clipping;
- (Random) extraction of fixed-size slices.

1.3 Deep network building

The building of a deep network is straightforward and can be done withing the same INI configuration file. Several layer types are available: convolutional, pooling, fully connected, Radial-basis function (RBF) and softmax. The tool is highly modular and new layer types can be added without any change in the code base. Parameters of each layer type are modifiable, for example for the convolutional layer, one can specify the size of the convolution kernels, the stride, the number of kernels per input map and the learning parameters (learning rate, initial weights value...). For the learning, the data dynamic can be chosen between 16 bits (with NVIDIA® cuDNN²), 32 bit and 64 bit floating point numbers.

The following example, which will serve as the use case for the rest of this presentation, shows how to build a DNN with 5 layers: one convolution layer, followed by one MAX pooling layer, followed by two fully connected layers and a softmax output layer.

```

; Specify the input data format
[env]
SizeX=24
SizeY=24
BatchSize=12

; First layer: convolutional with 3x3 kernels
[conv1]
Input=env
Type=Conv

```

²On future GPUs

```

KernelWidth=3
KernelHeight=3
NbOutputs=32
Stride=1

; Second layer: MAX pooling with pooling area 2x2
[pool1]
Input=conv1
Type=Pool
Pooling=Max
PoolWidth=2
PoolHeight=2
NbOutputs=32
Stride=2
Mapping.Size=1 ; one to one connection between convolution output maps and pooling input maps

; Third layer: fully connected layer with 60 neurons
[fc1]
Input=pool1
Type=Fc
NbOutputs=60

; Fourth layer: fully connected with 10 neurons
[fc2]
Input=fc1
Type=Fc
NbOutputs=10

; Final layer: softmax
[softmax]
Input=fc2
Type=Softmax
NbOutputs=10
WithLoss=1

[softmax.Target]
TargetValue=1.0
DefaultValue=0.0

```

The resulting DNN is shown in figure 1.

The learning is accelerated in GPU using the NVIDIA® cuDNN framework, integrated into the toolflow. Using GPU acceleration, learning times can be reduced typically by two orders of magnitude, enabling the learning of large databases within tens of minutes to a few hours instead of several days or weeks for non-GPU accelerated learning.

1.4 Performances evaluation

The software automatically outputs all the information needed for the network applicative performances analysis, such as the recognition rate and the validation score during the learning; the confusion matrix during learning, validation and test; the memory and computation requirements of the network; the output maps activity for each layer, and so on, as shown in figure 2.

1.5 Hardware exports

Once the learned DNN recognition rate performances are satisfying, an optimized version of the network can be automatically exported for various embedded targets. An automated network computation performances benchmarking can also be performed among different targets.

The following targets are currently supported by the toolflow:

- Plain C code (no dynamic memory allocation, no floating point processing);



Figure 1: Automatically generated and ready to learn DNN from the INI configuration file example.




Figure 2: Example of information automatically generated by the software during and after learning.

- C code accelerated with OpenMP;
- C code tailored for High-Level Synthesis (HLS) with Xilinx® Vivado® HLS;
Direct synthesis to FPGA, with timing and utilization after routing;

Possibility to constrain the maximum number of clock cycles desired to compute the whole network;

FPGA utilization vs number of clock cycle trade-off analysis;

- OpenCL code optimized for either CPU/DSP or GPU;
- Cuda kernels, cuDNN and TensorRT code optimized for NVIDIA® GPUs.

Different automated optimizations are embedded in the exports:

- DNN weights and signal data precision reduction (down to 8 bit integers or less for custom FPGA IPs);
- Non-linear network activation functions approximations;
- Different weights discretization methods.

The exports are generated automatically and come with a Makefile and a working testbench, including the pre-processed testing dataset. Once generated, the testbench is ready to be compiled and executed on the target platform. The applicative performance (recognition rate) as well as the computing time per input data can then be directly measured by the testbench.




Figure 3: Example of network benchmarking on different hardware targets.

The figure 3 shows an example of benchmarking results of the previous DNN on different targets (in log scale). Compared to desktop CPUs, the number of input image pixels processed per second is more than one order of magnitude higher with GPUs and at least two orders of magnitude better with synthesized DNN on FPGA.

1.6 Summary

The N2D2 platform is today a complete and production ready neural network building tool, which does not require advanced knowledges in deep learning to be used. It is tailored for fast neural network applications generation and porting with minimum overhead in terms of database creation and management, data pre-processing, networks configuration and optimized code generation, which can save months of manual porting and verification effort to a single automated step in the tool.

2 About N2D2-IP

While N2D2 is our deep learning open-source core framework, some modules referred as "N2D2-IP" in the manual, are only available through custom license agreement with CEA LIST.

If you are interested in obtaining some of these modules, please contact our business developer for more information on available licensing options:

Sandrine VARENNE (Sandrine.VARENNE@cea.fr)

In addition to N2D2-IP modules, we can also provide our expertise to design specific solutions for integrating DNN in embedded hardware systems, where power, latency, form factor and/or cost are constrained. We can target CPU/DSP/GPU CoTS hardware as well as our own PNeuro (programmable) and DNeuro (dataflow) dedicated hardware accelerator IPs for DNN on FPGA or ASIC.

3 Performing simulations

3.1 Obtaining the latest version of this manual

Before going further, please make sure you are reading the latest version of this manual. It is located in the `manual` sub-directory. To compile the manual in PDF, just run the following command:

```
cd manual && make
```

In order to compile the manual, you must have `pdflatex` and `bibtex` installed, as well as some common LaTeX packages.

- On Ubuntu, this can be done by installing the `texlive` and `texlive-latex-extra` software packages.
- On Windows, you can install the `miKTeX` software, which includes everything needed and will install the required LaTeX packages on the fly.

3.2 Minimum system requirements

- Supported processors:
 - ARM Cortex A15 (tested on Tegra K1)
 - ARM Cortex A53/A57 (tested on Tegra X1)
 - Pentium-compatible PC (Pentium III, Athlon or more-recent system recommended)
- Supported operating systems:
 - Windows ≥ 7 or Windows Server ≥ 2012 , 64 bits with Visual Studio ≥ 2015.2 (2015 Update 2)
 - GNU/Linux with GCC ≥ 4.4 (tested on RHEL ≥ 6 , Debian ≥ 6 , Ubuntu ≥ 14.04)
- At least 256 MB of RAM (1 GB with GPU/CUDA) for MNIST dataset processing
- At least 150 MB available hard disk space + 350 MB for MNIST dataset processing

For CUDA acceleration:

- CUDA ≥ 6.5 and CuDNN ≥ 1.0
- NVIDIA GPU with CUDA compute capability ≥ 3 (starting from *Kepler* micro-architecture)
- At least 512 MB GPU RAM for MNIST dataset processing

3.3 Obtaining N2D2

3.3.1 Prerequisites

Red Hat Enterprise Linux (RHEL) 6 Make sure you have the following packages installed:

- `cmake`
- `gnuplot`
- `opencv`
- `opencv-devel` (may require the `rhel-x86_64-workstation-optional-6` repository channel)

Plus, to be able to use GPU acceleration:

- Install the CUDA repository package:

```
rpm -Uhv http://developer.download.nvidia.com/compute/cuda/repos/rhel6/x86_64/cuda-repo-rhel6-7.5-18.x86_64.rpm
yum clean expire-cache
yum install cuda
```
- Install cuDNN from the NVIDIA website: register to [NVIDIA Developer](#) and download the latest version of cuDNN. Simply copy the header and library files from the cuDNN archive to the corresponding directories in the CUDA installation path (by default: `/usr/local/cuda/include` and `/usr/local/cuda/lib64`, respectively).
- Make sure the CUDA library path (e.g. `/usr/local/cuda/lib64`) is added to the `LD_LIBRARY_PATH` environment variable.

Ubuntu Make sure you have the following packages installed, if they are available on your Ubuntu version:

- `cmake`
- `gnuplot`
- `libopencv-dev`
- `libcv-dev`
- `libhighgui-dev`

Plus, to be able to use GPU acceleration:

- Install the CUDA repository package matching your distribution. For example, for Ubuntu 14.04 64 bits:

```
wget http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1404/x86_64/cuda-repo-ubuntu1404_7.5-18_amd64.deb
dpkg -i cuda-repo-ubuntu1404_7.5-18_amd64.deb
```

- Install the cuDNN repository package matching your distribution. For example, for Ubuntu 14.04 64 bits:

```
wget http://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1404/x86_64/nvidia-machine-learning-repo-ubuntu1404_4.0-2_amd64.deb
dpkg -i nvidia-machine-learning-repo-ubuntu1404_4.0-2_amd64.deb
```

Note that the cuDNN repository package is provided by NVIDIA for Ubuntu starting from version 14.04.

- Update the package lists: `apt-get update`
- Install the CUDA and cuDNN required packages:

```
apt-get install cuda-core-7-5 cuda-cudart-dev-7-5 cuda-cublas-dev-7-5 cuda-curand-dev-7-5 libcudnn5-dev
```

- Make sure there is a symlink to `/usr/local/cuda`:

```
ln -s /usr/local/cuda-7.5 /usr/local/cuda
```
- Make sure the CUDA library path (e.g. `/usr/local/cuda/lib64`) is added to the `LD_LIBRARY_PATH` environment variable.

Windows On Windows 64 bits, Visual Studio \geq 2015.2 (2015 Update 2) is required.

Make sure you have the following software installed:

- CMake (<http://www.cmake.org/>): download and run the Windows installer.
- `dirent.h` C++ header (<https://github.com/tronkko/dirent>): to be put in the Visual Studio include path.
- Gnuplot (<http://www.gnuplot.info/>): the bin sub-directory in the install path needs to be added to the Windows PATH environment variable.
- OpenCV (<http://opencv.org/>): download the latest 2.x version for Windows and extract it to, for example, `C:\OpenCV\`. Make sure to define the environment variable `OpenCV_DIR` to point to `C:\OpenCV\opencv\build`. Make sure to add the bin sub-directory (`C:\OpenCV\opencv\build\x64\vc12\bin`) to the Windows PATH environment variable.

Plus, to be able to use GPU acceleration:

- Download and install CUDA toolkit 8.0 located at https://developer.nvidia.com/compute/cuda/8.0/prod/local_installers/cuda_8.0.44_windows-exe:

```
rename cuda_8.0.44_windows-exe cuda_8.0.44_windows.exe
cuda_8.0.44_windows.exe -s compiler_8.0 cublas_8.0 cublas_dev_8.0 cudart_8.0 curand_8.0
                        curand_dev_8.0
```

- Update the PATH environment variable:

```
set PATH=%ProgramFiles%\NVIDIA GPU Computing Toolkit\CUDA\v8.0\bin;%ProgramFiles%\NVIDIA GPU
      Computing Toolkit\CUDA\v8.0\libnvvp;%PATH%
```

- Download and install cuDNN 8.0 located at <http://developer.download.nvidia.com/compute/redist/cudnn/v5.1/cudnn-8.0-windows7-x64-v5.1.zip> (the following command assumes that you have 7-Zip installed):

```
7z x cudnn-8.0-windows7-x64-v5.1.zip
copy cuda\include\*. * ^
    "%ProgramFiles%\NVIDIA GPU Computing Toolkit\CUDA\v8.0\include\"
copy cuda\lib\x64\*. * ^
    "%ProgramFiles%\NVIDIA GPU Computing Toolkit\CUDA\v8.0\lib\x64\"
copy cuda\bin\*. * ^
    "%ProgramFiles%\NVIDIA GPU Computing Toolkit\CUDA\v8.0\bin\"
```

3.3.2 Getting the sources

Use the following command:

```
git clone git@github.com:CEA-LIST/N2D2.git
```

3.3.3 Compilation

To compile the program:

```
mkdir build
cd build
cmake .. && make
```

On Windows, you may have to specify the generator, for example:

```
cmake .. -G"Visual Studio 14"
```

Then open the newly created N2D2 project in Visual Studio 2015. Select "Release" for the build target. Right click on `ALL_BUILD` item and select "Build".

3.4 Downloading training datasets

A python script located in the repository root directory allows you to select and automatically download some well-known datasets, like MNIST and GTSRB (the script requires Python 2.x with bindings for GTK 2 package):

```
./tools/install_stimuli_gui.py
```

By default, the datasets are downloaded in the path specified in the `N2D2_DATA` environment variable, which is the root path used by the N2D2 tool to locate the databases. If the `N2D2_DATA` variable is not set, the default value used is `/local/$USER/n2d2_data/` (or `/local/n2d2_data/` if the `USER` environment variable is not set) on Linux and `C:\n2d2_data\` on Windows.

Please make sure you have write access to the `N2D2_DATA` path, or if not set, in the default `/local/$USER/n2d2_data/` path.

3.5 Run the learning

The following command will run the learning for 600,000 image presentations/steps and log the performances of the network every 10,000 steps:

```
./n2d2 "mnist24_16c4s2_24c5s2_150_10.ini" -learn 600000 -log 10000
```

Note: you may want to check the gradient computation using the `-check` option. Note that it can be extremely long and can occasionally fail if the required precision is too high.

3.6 Test a learned network

After the learning is completed, this command evaluate the network performances on the test data set:

```
./n2d2 "mnist24_16c4s2_24c5s2_150_10.ini" -test
```

3.6.1 Interpreting the results

Recognition rate The recognition rate and the validation score are reported during the learning in the `TargetScore_*/Success_validation.png` file, as shown in figure 4.

Confusion matrix The software automatically outputs the confusion matrix during learning, validation and test, with an example shown in figure 5. Each row of the matrix contains the number of occurrences estimated by the network for each label, for all the data corresponding to a single actual, target label. Or equivalently, each column of the matrix contains the number of actual, target label occurrences, corresponding to the same estimated label. Ideally, the matrix should be diagonal, with no occurrence of an estimated label for a different actual label (network mistake).

The confusion matrix reports can be found in the simulation directory:

- `TargetScore_*/ConfusionMatrix_learning.png`;
- `TargetScore_*/ConfusionMatrix_validation.png`;
- `TargetScore_*/ConfusionMatrix_test.png`.

Memory and computation requirements The software also report the memory and computation requirements of the network, as shown in figure 6. The corresponding report can be found in the `stats` sub-directory of the simulation.

Kernels and weights distribution The synaptic weights obtained during and after the learning can be analyzed, in terms of distribution (`weights` sub-directory of the simulation) or in terms of kernels (`kernels` sub-directory of the simulation), as shown in 7.




Figure 4: Recognition rate and validation score during learning.




Figure 5: Example of confusion matrix obtained after the learning.

Output maps activity The initial output maps activity for each layer can be visualized in the *outputs_init* sub-directory of the simulation, as shown in figure 8.

3.7 Export a learned network



Figure 6: Example of memory and computation requirements of the network.

```
./n2d2 "mnist24_16c4s2_24c5s2_150_10.ini" -export CPP_OpenCL
```

Export types:

- `c` C export using OpenMP;
- `c_HLS` C export tailored for HLS with Vivado HLS;
- `CPP_OpenCL` C++ export using OpenCL;
- `CPP_Cuda` C++ export using Cuda;
- `CPP_cuDNN` C++ export using cuDNN;
- `CPP_TensorRT` C++ export using tensorRT 2.1 API;
- `SC_Spike` SystemC spike export.

Other program options related to the exports:

Option [default value]	Description
<code>-nbbits</code> [8]	Number of bits for the weights and signals. Must be 8, 16, 32 or 64 for integer export, or -32, -64 for floating point export. The number of bits can be arbitrary for the <code>c_HLS</code> export (for example, 6 bits). It must be -32 for the <code>CPP_TensorRT</code> export, the precision is directly set at runtime
<code>-calib</code> [0]	Number of stimuli used for the calibration. 0 = no calibration (default), -1 = use the full test dataset for calibration
<code>-calib-passes</code> [2]	Number of KL passes for determining the layer output values distribution truncation threshold (0 = use the max. value, no truncation)
<code>-no-unsigned</code>	If present, disable the use of unsigned data type in integer exports
<code>-db-export</code> [-1]	Max. number of stimuli to export (0 = no dataset export, -1 = unlimited)




Figure 7: Example of kernels and weights distribution analysis for two convolutional layers.

3.7.1 C export

Test the exported network:

```
cd export_C_int8
make
./bin/n2d2_test
```

The result should look like:

```
...
1652.00/1762      (avg = 93.757094%)
1653.00/1763      (avg = 93.760635%)
1654.00/1764      (avg = 93.764172%)
Tested 1764 stimuli
Success rate = 93.764172%
Process time per stimulus = 187.548186 us (12 threads)
```

Confusion matrix:

/ T \ E /	0 /	1 /	2 /	3 /



Figure 8: Output maps activity example of the first convolutional layer of the network.

/	0	/	329	/	1	/	5	/	2	/
/		/	97.63%	/	0.30%	/	1.48%	/	0.59%	/
/	1	/	0	/	692	/	2	/	6	/
/		/	0.00%	/	98.86%	/	0.29%	/	0.86%	/
/	2	/	11	/	27	/	609	/	55	/
/		/	1.57%	/	3.85%	/	86.75%	/	7.83%	/
/	3	/	0	/	0	/	1	/	24	/
/		/	0.00%	/	0.00%	/	4.00%	/	96.00%	/

T: Target E: Estimated

3.7.2 CPP_OpenCL export

The OpenCL export can run the generated program in GPU or CPU architectures. Compilation features:

Preprocessor command [default value]	Description
PROFILING [0]	Compile the binary with a synchronization between each layers and return the mean execution time of each layer. This preprocessor option can decrease performances.
GENERATE_KBIN [0]	Generate the binary output of the OpenCL kernel .cl file use. The binary is store in the /bin folder.
LOAD_KBIN [0]	Indicate to the program to load an OpenCL kernel as a binary from the /bin folder instead of a .cl file.
CUDA [0]	Use the CUDA OpenCL SDK locate at <i>/usr/local/cuda</i>
MALI [0]	Use the MALI OpenCL SDK locate at <i>/usr/MaliOpenCLSDKvXXX</i>
INTEL [0]	Use the INTEL OpenCL SDK locate at <i>/opt/intel/opencl</i>
AMD [1]	Use the AMD OpenCL SDK locate at <i>/opt/AMDAPPSDK - XXX</i>

Program options related to the OpenCL export:

Option [default value]	Description
-cpu	If present, force to use a CPU architecture to run the program
-gpu	If present, force to use a GPU architecture to run the program
-batch [1]	Size of the batch to use
-stimulus [NULL]	Path to a specific input stimulus to test. For example: -stimulus <i>/stimulus/env0000.pgm</i> command will test the file env0000.pgm of the stimulus folder.

Test the exported network:

```
cd export_CPP_OpenCL_float32
make
./bin/n2d2_openc1_test -gpu
```

3.7.3 CPP_TensorRT export

The TensorRT API export can run the generated program in NVIDIA GPU architecture. It use CUDA, cuDNN and TensorRT API library. All the native TensorRT layers are supported. The export support from TensorRT 2.1 to TensorRT 5.0 versions.

Program options related to the TensorRT API export:

Option [default value]	Description
-batch [1]	Size of the batch to use
-dev [0]	CUDA Device ID selection
-stimulus [NULL]	Path to a specific input stimulus to test. For example: -stimulus /stimulus/env0000.pgm command will test the file env0000.pgm of the stimulus folder.
-prof	Activates the layer wise profiling mechanism. This option can decrease execution time performance.
-iter-build [1]	Sets the number of minimization build iterations done by the tensorRT builder to find the best layer tactics.
-nbbits [-32]	Number of bits used for computation. Value -32 for Full FP32 bits configuration, -16 for Half FP16 bits configuration and 8 for INT8 bits configuration. When running INT8 mode for the first time, the TensorRT calibration process can be very long. Once generated the generated calibration table will be automatically reused. Supported compute mode in function of the compute capability are provided here: https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#compute-capabilities .

Test the exported network with layer wise profiling:

```
cd export_CPP_TensorRT_float32
make
./bin/n2d2_tensorRT_test -prof
```

The results of the layer wise profiling should look like:

```
(19%) ***** CONV1 + CONV1_ACTIVATION:
0.0219467 ms
(05%) ***** POOL1: 0.00675573 ms
(13%) ***** CONV2 + CONV2_ACTIVATION: 0.0159089 ms
(05%) ***** POOL2: 0.00616047 ms
(14%) ***** CONV3 + CONV3_ACTIVATION: 0.0159713 ms
(19%) ***** FC1 + FC1_ACTIVATION: 0.0222242 ms
(13%) ***** FC2: 0.0149013 ms
(08%) ***** SOFTMAX: 0.0100633 ms
Average profiled tensorRT process time per stimulus = 0.113932 ms
```

3.7.4 CPP_cuDNN export

The cuDNN export can run the generated program in NVIDIA GPU architecture. It use CUDA and cuDNN library. Compilation features:

Preprocessor command [default value]	Description
PROFILING [0]	Compile the binary with a synchronization between each layers and return the mean execution time of each layer. This preprocessor option can decrease performances.
ARCH32 [0]	Compile the binary with the 32-bits architecture compatibility.

Program options related to the cuDNN export:

Option [default value]	Description
-batch [1] -dev [0] -stimulus [NULL]	Size of the batch to use CUDA Device ID selection Path to a specific input stimulus to test. For example: -stimulus / <i>stimulus/env0000.pgm</i> command will test the file env0000.pgm of the stimulus folder.

Test the exported network:

```
cd export_CPP_cuDNN_float32
make
./bin/n2d2_cudnn_test
```

3.7.5 C_HLS export

Test the exported network:

```
cd export_C_HLS_int8
make
./bin/n2d2_test
```

Run the High-Level Synthesis (HLS) with Xilinx® Vivado® HLS:

```
vivado_hls -f run_hls.tcl
```

3.7.6 Layer compatibility table

Layer compatibility table in function of the export type:

Layer compatibility table	Export Type			
	C	C_HLS	CPP_OpenCL	CPP_TensorRT
Conv	✓	✓	✓	✓
Pool	✓	✓	✓	✓
Fc	✓	✓	✓	✓
Softmax	✓	✗	✓	✓
FMP	✓	✗	✓	✗
Deconv	✗	✗	✗	✓
ElemWise	✗	✗	✗	✓
Resize	✓	✗	✗	✓
Padding	✗	✗	✗	✓
LRN	✗	✗	✗	✓
Anchor	✗	✗	✗	✓
ObjectDet	✗	✗	✗	✓
ROIpooling	✗	✗	✗	✓
RP	✗	✗	✗	✓

BatchNorm is not mentionned because batch normalization parameters are automatically fused with convolutions parameters with the command "-fuse".

4 INI file interface

The INI file interface is the primary way of using N2D2. It is a simple, lightweight and user-friendly format for specifying a complete DNN-based application, including dataset instantiation, data pre-processing, neural network layers instantiation and post-processing, with all its hyperparameters.

4.1 Syntax

INI files are simple text files with a basic structure composed of sections, properties and values.

4.1.1 Properties

The basic element contained in an INI file is the property. Every property has a name and a value, delimited by an equals sign (=). The name appears to the left of the equals sign.

```
name=value
```

4.1.2 Sections

Properties may be grouped into arbitrarily named sections. The section name appears on a line by itself, in square brackets ([and]). All properties after the section declaration are associated with that section. There is no explicit "end of section" delimiter; sections end at the next section declaration, or the end of the file. Sections may not be nested.

```
[section]  
a=a  
b=b
```

4.1.3 Case sensitivity

Section and property names are case sensitive.

4.1.4 Comments

Semicolons (;) or number sign (#) at the beginning or in the middle of the line indicate a comment. Comments are ignored.

```
; comment text  
a=a # comment text  
a="a ; not a comment" ; comment text
```

4.1.5 Quoted values

Values can be quoted, using double quotes. This allows for explicit declaration of whitespace, and/or for quoting of special characters (equals, semicolon, etc.).

4.1.6 Whitespace

Leading and trailing whitespace on a line are ignored.

4.1.7 Escape characters

A backslash (\) followed immediately by EOL (end-of-line) causes the line break to be ignored.

4.2 Template inclusion syntax

Is is possible to recursively include templated INI files. For example, the main INI file can include a templated file like the following:

```
[inception@inception_model.ini.tpl]
INPUT=layer_x
SIZE=32
ARRAY=2 ; Must be the number of elements in the array
ARRAY[0].P1=Conv
ARRAY[0].P2=32
ARRAY[1].P1=Pool
ARRAY[1].P2=64
```

If the inception_model.ini.tpl template file content is:

```
[{{SECTION_NAME}}_layer1]
Input={{INPUT}}
Type=Conv
NbOutputs={{SIZE}}

[{{SECTION_NAME}}_layer2]
Input={{SECTION_NAME}}_layer1
Type=Fc
NbOutputs={{SIZE}}

{% block ARRAY %}
[{{SECTION_NAME}}_array{{#}}]
Prop1=Config{{.P1}}
Prop2={{.P2}}
{% endblock %}
```

The resulting equivalent content for the main INI file will be:

```
[inception_layer1]
Input=layer_x
Type=Conv
NbOutputs=32

[inception_layer2]
Input=inception_layer1
Type=Fc
NbOutputs=32

[inception_array0]
Prop1=ConfigConv
Prop2=32

[inception_array1]
Prop1=ConfigPool
Prop2=64
```

The SECTION_NAME template parameter is automatically generated from the name of the including section (before @).

4.2.1 Variable substitution

{{VAR}} is replaced by the value of the VAR template parameter.

4.2.2 Control statements

Control statements are between {% and %} delimiters.

block `{%block ARRAY %} ... {%endblock %}`

The `#` template parameter is automatically generated from the `{%block ... %}` template control statement and corresponds to the current item position, starting from 0.

for `{%for VAR in range([START,]END])%} ... {%endfor %}`

If `START` is not specified, the loop begins at 0 (first value of `VAR`). The last value of `VAR` is `END-1`.

if `{%if VAR OP [VALUE] %} ... [{%else %}] ... {%endif %}`

`OP` may be `==`, `!=`, `exists` OR `not_exists`.

include `{%include FILENAME %}`

4.3 Global parameters

Option [default value]	Description
DefaultModel [Transcode]	Default layers model. Can be <code>Frame</code> , <code>Frame_CUDA</code> , <code>Transcode</code> OR <code>Spike</code>
DefaultDataType [Float32]	Default layers data type. Can be <code>Float16</code> , <code>Float32</code> OR <code>Float64</code>
SignalsDiscretization [0]	Number of levels for signal discretization
FreeParametersDiscretization [0]	Number of levels for weights discretization

4.4 Databases

The tool integrates pre-defined modules for several well-known database used in the deep learning community, such as MNIST, GTSRB, CIFAR10 and so on. That way, no extra step is necessary to be able to directly build a network and learn it on these database.

4.4.1 MNIST

MNIST ([LeCun et al., 1998](#)) is already fractionned into a learning set and a testing set, with:

- 60,000 digits in the learning set;
- 10,000 digits in the testing set.

Example:

```
[database]
Type=MNIST_IDX_Database
Validation=0.2 ; Fraction of learning stimuli used for the validation [default: 0.0]
```

Option [default value]	Description
Validation [0.0]	Fraction of the learning set used for validation
DataPath [\$N2D2_DATA/mnist]	Path to the database

4.4.2 GTSRB

GTSRB ([Stallkamp et al., 2012](#)) is already fractionned into a learning set and a testing set, with:

- 39,209 digits in the learning set;
- 12,630 digits in the testing set.

Example:

[database]

Type=GTSRB_DIR_Database

Validation=0.2 ; Fraction of learning stimuli used for the validation [default: 0.0]

Option [default value]	Description
Validation [0.0]	Fraction of the learning set used for validation
DataPath [\$N2D2_DATA/GTSRB]	Path to the database

4.4.3 Directory

Hand made database stored in files directories are directly supported with the `DIR_Database` module. For example, suppose your database is organized as following (in the path specified in the `N2D2_DATA` environment variable):

- GST/airplanes: 800 images
- GST/car_side: 123 images
- GST/Faces: 435 images
- GST/Motorbikes: 798 images

You can then instantiate this database as input of your neural network using the following parameters:

[database]

Type=DIR_Database

DataPath=\${N2D2_DATA}/GST

Learn=0.4 ; 40% of images of the smallest category = 49 (0.4x123) images for each category will be used for learning

Validation=0.2 ; 20% of images of the smallest category = 25 (0.2x123) images for each category will be used for validation

; the remaining images will be used for testing

Each subdirectory will be treated as a different label, so there will be 4 different labels, named after the directory name.

The stimuli are equi-partitioned for the learning set and the validation set, meaning that the same number of stimuli for each category is used. If the learn fraction is 0.4 and the validation fraction is 0.2, as in the example above, the partitioning will be the following:

Label ID	Label name	Learn set	Validation set	Test set
0	airplanes	49	25	726
1	car_side	49	25	49
2	Faces	49	25	361
3	Motorbikes	49	25	724
Total:		196	100	1860

 *Mandatory option*

Option [default value]	Description
DataPath	Path to the root stimuli directory
Learn	If <code>PerLabelPartitioning</code> is true, fraction of images used for the learning; else, number of images used for the learning, regardless of their labels
LoadInMemory [0]	Load the whole database into memory
Depth [1]	Number of sub-directory levels to include. Examples:

<p>LabelName []</p> <p>LabelDepth [1]</p> <p>PerLabelPartitioning [1]</p> <p>Validation [0.0]</p> <p>Test [1.0-Learn-Validation]</p> <p>ValidExtensions []</p> <p>LoadMore []</p>	<p>Depth = 0: load stimuli only from the current directory (DataPath)</p> <p>Depth = 1: load stimuli from DataPath and stimuli contained in the sub-directories of DataPath</p> <p>Depth < 0: load stimuli recursively from DataPath and all its sub-directories</p> <p>Base stimuli label name</p> <p>Number of sub-directory name levels used to form the stimuli labels. Examples:</p> <p>LabelDepth = -1: no label for all stimuli (label ID = -1)</p> <p>LabelDepth = 0: uses LabelName for all stimuli</p> <p>LabelDepth = 1: uses LabelName for stimuli in the current directory (DataPath) and LabelName/<i>sub-directory name</i> for stimuli in the sub-directories</p> <p>If true, the stimuli are equi-partitioned for the learn/validation/test sets, meaning that the same number of stimuli for each label is used</p> <p>If PerLabelPartitioning is true, fraction of images used for the validation; else, number of images used for the validation, regardless of their labels</p> <p>If PerLabelPartitioning is true, fraction of images used for the test; else, number of images used for the test, regardless of their labels</p> <p>List of space-separated valid stimulus file extensions (if left empty, any file extension is considered a valid stimulus)</p> <p>Name of an other section with the same options to load a different DataPath</p>
<p>ROIFile []</p> <p>DefaultLabel []</p> <p>ROIsMargin [0]</p>	<p>File containing the stimuli ROIs. If a ROI file is specified, LabelDepth should be set to -1</p> <p>Label name for pixels outside any ROI (default is no label, pixels are ignored)</p> <p>Number of pixels around ROIs that are ignored (and not considered as DefaultLabel pixels)</p>

To load and partition more than one DataPath, one can use the LoadMore option:

```
[database]
Type=DIR_Database
DataPath=${N2D2_DATA}/GST
Learn=0.6
Validation=0.4
LoadMore=database.test

; Load stimuli from the "GST_Test" path in the test dataset
[database.test]
DataPath=${N2D2_DATA}/GST_Test
Learn=0.0
Test=1.0
; The LoadMore option is recursive:
; LoadMore=database.more

; [database.more]
; Load even more data here
```

Speech Commands Dataset Use with Speech Commands Data Set, released by the Google (Warden, 2018).

```
[database]
Type=DIR_Database
DataPath=${N2D2_DATA}/speech_commands_v0.02
ValidExtensions=wav
IgnoreMasks=*/_background_noise_
Learn=0.6
Validation=0.2
```

4.4.4 Other built-in databases

Actitracker_Database Actitracker database, released by the WISDM Lab (W. Lockhart et al., 2011).

Option [default value]	Description
Learn [0.6]	Fraction of data used for the learning
Validation [0.2]	Fraction of data used for the validation
UseUnlabeledForTest [0]	If true, use the unlabeled dataset for the test
DataPath [<code>\${N2D2_DATA}</code> / <code>WISDM_at_v2.0</code>]	Path to the database

CIFAR10_Database CIFAR10 database (Krizhevsky, 2009).

Option [default value]	Description
Validation [0.0]	Fraction of the learning set used for validation
DataPath [<code>\${N2D2_DATA}/cifar-10-batches-bin</code>]	Path to the database

CIFAR100_Database CIFAR100 database (Krizhevsky, 2009).

Option [default value]	Description
Validation [0.0]	Fraction of the learning set used for validation
UseCoarse [0]	If true, use the coarse labeling (10 labels instead of 100)
DataPath [<code>\${N2D2_DATA}/cifar-100-binary</code>]	Path to the database

CKP_Database The Extended Cohn-Kanade (CK+) database for expression recognition (Lucey et al., 2010).

Option [default value]	Description
Learn	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
DataPath [<code>\${N2D2_DATA}/cohn-kanade-images</code>]	Path to the database

Caltech101_DIR_Database Caltech 101 database ([Fei-Fei et al., 2004](#)).

Option [default value]	Description
Learn	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
IncClutter [0]	If true, includes the BACKGROUND_Google directory of the database
DataPath [\$N2D2_DATA/ 101_ObjectCategories]	Path to the database

Caltech256_DIR_Database Caltech 256 database ([Griffin et al., 2007](#)).

Option [default value]	Description
Learn	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
IncClutter [0]	If true, includes the BACKGROUND_Google directory of the database
DataPath [\$N2D2_DATA/ 256_ObjectCategories]	Path to the database

CaltechPedestrian_Database Caltech Pedestrian database ([Dollár et al., 2009](#)).

Note that the images and annotations must first be extracted from the seq video data located in the *videos* directory using the `dbExtract.m` Matlab tool provided in the "Matlab evaluation/labeling code" downloadable on the dataset website.

Assuming the following directory structure (in the path specified in the `N2D2_DATA` environment variable):

- CaltechPedestrians/data-USA/videos/... (from the *setxx.tar* files)
- CaltechPedestrians/data-USA/annotations/... (from the *setxx.tar* files)
- CaltechPedestrians/tools/piotr_toolbox/toolbox (from the Piotr's Matlab Toolbox archive)
- CaltechPedestrians/*.m including `dbExtract.m` (from the Matlab evaluation/labeling code)

Use the following command in Matlab to generate the images and annotations:

```
cd([getenv('N2D2_DATA') '/CaltechPedestrians'])
addpath(genpath('tools/piotr_toolbox/toolbox')) % add the Piotr's Matlab Toolbox in the Matlab
path
dbInfo('USA')
dbExtract()
```

Option [default value]	Description
Validation [0.0]	Fraction of the learning set used for validation
SingleLabel [1]	Use the same label for "person" and "people" bounding box
IncAmbiguous [0]	Include ambiguous bounding box labeled "person?" using the same label as "person"
DataPath [\$N2D2_DATA/ CaltechPedestrians/data- USA/images]	Path to the database images
LabelPath [\$N2D2_DATA/ CaltechPedestrians/data- USA/annotations]	Path to the database annotations

Cityscapes_Database Cityscapes database (Cordts et al., 2016).

Option [default value]	Description
IncTrainExtra [0]	If true, includes the left 8-bit images - trainextra set (19,998 images)
UseCoarse [0]	If true, only use coarse annotations (which are the only annotations available for the trainextra set)
SingleInstanceLabels [1]	If true, convert group labels to single instance labels (for example, <code>cargroup</code> becomes <code>car</code>)
DataPath [\$N2D2_DATA/ Cityscapes/leftImg8bit] or [\$CITYSCAPES_DATASET] if defined	Path to the database images
LabelPath []	Path to the database annotations (deduced from <code>DataPath</code> if left empty)

Daimler_Database Daimler Monocular Pedestrian Detection Benchmark (Daimler Pedestrian).

Option [default value]	Description
Learn [1.0]	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
Test [0.0]	Fraction of images used for the test
Fully [0]	When activate it use the test dataset to learn. Use only on fully-cnn mode

DOTA_Database DOTA database (Xia et al., 2017).

Option [default value]	Description
Learn	Fraction of images used for the learning
DataPath [\$N2D2_DATA/DOTA]	Path to the database
LabelPath []	Path to the database labels list file

FDDB_Database Face Detection Data Set and Benchmark (FDDB) (Jain and Learned-Miller, 2010).

Option [default value]	Description
Learn	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
DataPath [\$N2D2_DATA/FDDB]	Path to the images (decompressed originalPics.tar.gz)
LabelPath [\$N2D2_DATA/FDDB]	Path to the annotations (decompressed FDDB-folds.tgz)

GTSDb_DIR_Database GTSDb database (Houben et al., 2013).

Option [default value]	Description
Learn	Fraction of images used for the learning
Validation [0.0]	Fraction of images used for the validation
DataPath [\$N2D2_DATA/FullIJCNN2013]	Path to the database

ILSVRC2012_Database ILSVRC2012 database (Russakovsky et al., 2015).

Option [default value]	Description
Learn	Fraction of images used for the learning
DataPath [\$N2D2_DATA/ILSVRC2012]	Path to the database
LabelPath [\$N2D2_DATA /ILSVRC2012/synsets.txt]	Path to the database labels list file

KITTI_Database The KITTI Database provide ROI which can be use for autonomous driving and environment perception. The database provide 8 labeled different classes. Utilization of the KITTI Database is under licensing conditions and request an email registration. To install it you have to follow this link: http://www.cvlibs.net/datasets/kitti/eval_tracking.php and download the left color images (15 GB) and the training labels of tracking data set (9 MB). Extract the downloaded archives in your **\$N2D2_DATA/KITTI** folder.

Option [default value]	Description
Learn [0.8]	Fraction of images used for the learning
Validation [0.2]	Fraction of images used for the validation

KITTI_Road_Database The KITTI Road Database provide ROI which can be used to road segmentation. The dataset provide 1 labeled class (road) on 289 training images. The 290 test images are not labeled. Utilization of the KITTI Road Database is under licensing conditions and request an email registration. To install it you have to follow this link: http://www.cvlibs.net/datasets/kitti/eval_road.php and download the "base kit" of (0.5 GB) with left color images, calibration and training labels. Extract the downloaded archive in your **\$N2D2_DATA/KITTI** folder.

Option [default value]	Description
Learn [0.8]	Fraction of images used for the learning
Validation [0.2]	Fraction of images used for the validation

KITTI_Object_Database The KITTI Object Database provide ROI which can be use for autonomous driving and environment perception. The database provide 8 labeled different classes on 7481 training images. The 7518 test images are not labeled. The whole database provide 80256 labeled objects. Utilization of the KITTI Object Database is under licensing conditions and request an email registration. To install it you have to follow this link: http://www.cvlibs.net/datasets/kitti/eval_object.php?obj_benchmark and download the "left color images" (12 GB) and the training labels of object data set (5 MB). Extract the downloaded archives in your **\$N2D2_DATA/KITTI_Object** folder.

Option [default value]	Description
Learn [0.8]	Fraction of images used for the learning
Validation [0.2]	Fraction of images used for the validation

LITISRouen_Database LITIS Rouen audio scene dataset (Rakotomamonjy and Gasso, 2014).

Option [default value]	Description
Learn [0.4]	Fraction of images used for the learning
Validation [0.4]	Fraction of images used for the validation
DataPath [\$N2D2_DATA/data_rouen]	Path to the database

4.4.5 Dataset images slicing

It is possible to automatically slice images from a dataset, with a given slice size and stride, using the `.slicing` attribute. This effectively increases the number of stimuli in the set.

```
[database.slicing]
ApplyTo=NoLearn
Width=2048
Height=1024
StrideX=2048
StrideY=1024
RandomShuffle=1 ; 1 is the default value
```

The `RandomShuffle` option, enabled by default, randomly shuffle the dataset after slicing. If disabled, the slices are added in order at the end of the dataset.

4.5 Stimuli data analysis

You can enable stimuli data reporting with the following section (the name of the section must start with `env.StimuliData`):

```
[env.StimuliData-raw]
ApplyTo=LearnOnly
LogSizeRange=1
LogValueRange=1
```

The stimuli data reported for the full MNIST learning set will look like:

```
env.StimuliData-raw data:
Number of stimuli: 60000
Data width range: [28, 28]
Data height range: [28, 28]
Data channels range: [1, 1]
Value range: [0, 255]
Value mean: 33.3184
Value std. dev.: 78.5675
```

4.5.1 Zero-mean and unity standard deviation normalization

It is possible to normalize the whole database to have zero mean and unity standard deviation on the learning set using a `RangeAffineTransformation` transformation:

```
; Stimuli normalization based on learning set global mean and std.dev.
[env.Transformation-normalize]
Type=RangeAffineTransformation
FirstOperator=Minus
FirstValue=[env.StimuliData-raw]_GlobalValue.mean
SecondOperator=Divides
SecondValue=[env.StimuliData-raw]_GlobalValue.stdDev
```

The variables `_GlobalValue.mean` and `_GlobalValue.stdDev` are automatically generated in the `[env.StimuliData-raw]` block. Thanks to this facility, unknown and arbitrary database can be analysed and normalized in one single step without requiring any external data manipulation.

After normalization, the stimuli data reported is:

```
env.StimuliData-normalized data:
Number of stimuli: 60000
Data width range: [28, 28]
Data height range: [28, 28]
Data channels range: [1, 1]
Value range: [-0.424074, 2.82154]
Value mean: 2.64796e-07
Value std. dev.: 1
```

Where we can check that the global mean is close to 0 and the standard deviation is 1 on the whole dataset. The result of the transformation on the first images of the set can be checked in the generated *frames* folder, as shown in figure 9.




Figure 9: Image of the set after normalization.

4.5.2 Subtracting the mean image of the set

Using the `StimuliData` object followed with an `AffineTransformation`, it is also possible to use the mean image of the dataset to normalize the data:

```
[env.StimuliData-meanData]
ApplyTo=LearnOnly
MeanData=1 ; Provides the _MeanData parameter used in the transformation

[env.Transformation]
Type=AffineTransformation
FirstOperator=Minus
FirstValue=[env.StimuliData-meanData]_MeanData
```

The resulting global mean image can be visualized in `env.StimuliData-meanData/meanData.bin.png` and is shown in figure 10.

After this transformation, the reported stimuli data becomes:

```
env.StimuliData-processed data:
Number of stimuli: 60000
Data width range: [28, 28]
Data height range: [28, 28]
Data channels range: [1, 1]
Value range: [-139.554, 254.979]
Value mean: -3.45583e-08
Value std. dev.: 66.1288
```

The result of the transformation on the first images of the set can be checked in the generated *frames* folder, as shown in figure 11.




Figure 10: Global mean image generated by `StimuliData` with the `MeanData` parameter enabled.




Figure 11: Image of the set after the `AffineTransformation` subtracting the global mean image (keep in mind that the original image value range is $[0, 255]$).

4.6 Environment

The environment simply specify the input data format of the network (width, height and batch size). Example:

```
[env]
SizeX=24
SizeY=24
BatchSize=12 ; [default: 1]
```

Option [default value]	Description
SizeX	Environment width
SizeY	Environment height
NbChannels [1]	Number of channels (applicable only if there is no <code>env.ChannelTransformation[...]</code>)
BatchSize [1]	Batch size
CompositeStimuli [0]	If true, use pixel-wise stimuli labels
CachePath []	Stimuli cache path (no cache if left empty)
StimulusType [SingleBurst]	Method for converting stimuli into spike trains. Can be any of <code>SingleBurst</code> , <code>Periodic</code> , <code>JitteredPeriodic</code> or <code>Poissonian</code>
DiscardedLateStimuli [1.0]	The pixels in the pre-processed stimuli with a value above this limit never generate spiking events
PeriodMeanMin [50 TimeMs]	Mean minimum period $\overline{T_{min}}$, used for periodic temporal codings, corresponding to pixels in the pre-processed stimuli with a value of 0 (which are supposed to be the most significant pixels)
PeriodMeanMax [12 TimeS]	Mean maximum period $\overline{T_{max}}$, used for periodic temporal codings, corresponding to pixels in the pre-processed stimuli with a value of 1 (which are supposed to be the least significant pixels). This maximum period may be never reached if <code>DiscardedLateStimuli</code> is lower than 1.0
PeriodRelStdDev [0.1]	Relative standard deviation, used for periodic temporal codings, applied to the spiking period of a pixel
PeriodMin [11 TimeMs]	Absolute minimum period, or spiking interval, used for periodic temporal codings, for any pixel

4.6.1 Built-in transformations

There are 6 possible categories of transformations:

- `env.Transformation[...]` Transformations applied to the input images before channels creation;
- `env.OnTheFlyTransformation[...]` On-the-fly transformations applied to the input images before channels creation;
- `env.ChannelTransformation[...]` Create or add transformation for a specific channel;
- `env.ChannelOnTheFlyTransformation[...]` Create or add on-the-fly transformation for a specific channel;
- `env.ChannelsTransformation[...]` Transformations applied to all the channels of the input images;
- `env.ChannelsOnTheFlyTransformation[...]` On-the-fly transformations applied to all the channels of the input images.

Example:

```
[env.Transformation]
Type=PadCropTransformation
Width=24
Height=24
```

Several transformations can be applied successively. In this case, to be able to apply multiple transformations of the same category, a different suffix ([...]) must be added to each transformation.

The transformations will be processed in the order of appearance in the INI file regardless of their suffix.

Common set of parameters for any kind of transformation:

Option [default value]	Description
ApplyTo [All]	Apply the transformation only to the specified stimuli sets. Can be: LearnOnly: learning set only ValidationOnly: validation set only TestOnly: testing set only NoLearn: validation and testing sets only NoValidation: learning and testing sets only NoTest: learning and validation sets only All: all sets (default)

Example:

```
[env.Transformation-1]
Type=ChannelExtractionTransformation
CSChannel=Gray

[env.Transformation-2]
Type=RescaleTransformation
Width=29
Height=29

[env.Transformation-3]
Type=EqualizeTransformation

[env.OnTheFlyTransformation]
Type=DistortionTransformation
ApplyTo=LearnOnly ; Apply this transformation for the Learning set only
ElasticGaussianSize=21
ElasticSigma=6.0
ElasticScaling=20.0
Scaling=15.0
Rotation=15.0
```

List of available transformations:

AffineTransformation Apply an element-wise affine transformation to the image with matrixes of the same size.

Option [default value]	Description
FirstOperator	First element-wise operator, can be Plus, Minus, Multiplies, Divides
FirstValue	First matrix file name
SecondOperator [Plus]	Second element-wise operator, can be Plus, Minus, Multiplies, Divides
SecondValue []	Second matrix file name

The final operation is the following, with A the image matrix, B_{1st} , B_{2nd} the matrixes to add/subtract/multiply/divide and \odot the element-wise operator :

$$f(A) = (A \odot_{op_{1st}} B_{1st}) \odot_{op_{2nd}} B_{2nd}$$

ApodizationTransformation Apply an apodization window to each data row.

Option [default value]	Description
Size	Window total size (must match the number of data columns)
WindowName [Rectangular]	Window name. Possible values are: Rectangular : Rectangular Hann : Hann Hamming : Hamming Cosine : Cosine Gaussian : Gaussian Blackman : Blackman Kaiser : Kaiser

Gaussian window Gaussian window.

Option [default value]	Description
<i>WindowName.Sigma</i> [0.4]	Sigma

Blackman window Blackman window.

Option [default value]	Description
<i>WindowName.Alpha</i> [0.16]	Alpha

Kaiser window Kaiser window.

Option [default value]	Description
<i>WindowName.Beta</i> [5.0]	Beta

ChannelExtractionTransformation Extract an image channel.

Option	Description
CSChannel	Blue : blue channel in the BGR colorspace, or first channel of any colorspace Green : green channel in the BGR colorspace, or second channel of any colorspace Red : red channel in the BGR colorspace, or third channel of any colorspace Hue : hue channel in the HSV colorspace Saturation : saturation channel in the HSV colorspace Value : value channel in the HSV colorspace Gray : gray conversion y : Y channel in the YCbCr colorspace cb : Cb channel in the YCbCr colorspace cr : Cr channel in the YCbCr colorspace

ColorSpaceTransformation Change the current image colorspace.

Option	Description
ColorSpace	<p>BGR: convert any gray, BGR or BGRA image to BGR</p> <p>RGB: convert any gray, BGR or BGRA image to RGB</p> <p>HSV: convert BGR image to HSV</p> <p>HLS: convert BGR image to HLS</p> <p>YCrCb: convert BGR image to YCrCb</p> <p>CIELab: convert BGR image to CIELab</p> <p>CIEluv: convert BGR image to CIEluv</p> <p>RGB_to_BGR: convert RGB image to BGR</p> <p>RGB_to_HSV: convert RGB image to HSV</p> <p>RGB_to_HLS: convert RGB image to HLS</p> <p>RGB_to_YCrCb: convert RGB image to YCrCb</p> <p>RGB_to_CIELab: convert RGB image to CIELab</p> <p>RGB_to_CIEluv: convert RGB image to CIEluv</p> <p>HSV_to_BGR: convert HSV image to BGR</p> <p>HSV_to_RGB: convert HSV image to RGB</p> <p>HLS_to_BGR: convert HLS image to BGR</p> <p>HLS_to_RGB: convert HLS image to RGB</p> <p>YCrCb_to_BGR: convert YCrCb image to BGR</p> <p>YCrCb_to_RGB: convert YCrCb image to RGB</p> <p>CIELab_to_BGR: convert CIELab image to BGR</p> <p>CIELab_to_RGB: convert CIELab image to RGB</p> <p>CIEluv_to_BGR: convert CIEluv image to BGR</p> <p>CIEluv_to_RGB: convert CIEluv image to RGB</p>

Note that the default colorspace in N2D2 is BGR, the same as in OpenCV.

DFTTransformation Apply a DFT to the data. The input data must be single channel, the resulting data is two channels, the first for the real part and the second for the imaginary part.

Option [default value]	Description
TwoDimensional [1]	If true, compute a 2D image DFT. Otherwise, compute the 1D DFT of each data row

Note that this transformation can add zero-padding if required by the underlying FFT implementation.

DistortionTransformation Apply elastic distortion to the image. This transformation is generally used on-the-fly (so that a different distortion is performed for each image), and for the learning only.

Option [default value]	Description
ElasticGaussianSize [15]	Size of the gaussian for elastic distortion (in pixels)
ElasticSigma [6.0]	Sigma of the gaussian for elastic distortion
ElasticScaling [0.0]	Scaling of the gaussian for elastic distortion
Scaling [0.0]	Maximum random scaling amplitude (+/-, in percentage)
Rotation [0.0]	Maximum random rotation amplitude (+/-, in °)

EqualizeTransformation Image histogram equalization.

Option [default value]	Description
Method [Standard]	Standard: standard histogram equalization
CLAHE_ClipLimit [40.0]	CLAHE: contrast limited adaptive histogram equalization
CLAHE_GridSize [8]	Threshold for contrast limiting (for CLAHE only) Size of grid for histogram equalization (for CLAHE only). Input image will be divided into equally sized rectangular tiles. This parameter defines the number of tiles in row and column.

ExpandLabelTransformation Expand single image label (1x1 pixel) to full frame label.

FilterTransformation Apply a convolution filter to the image.

Option [default value]	Description
Kernel	Convolution kernel. Possible values are: *: custom kernel Gaussian: Gaussian kernel LoG: Laplacian Of Gaussian kernel DoG: Difference Of Gaussian kernel Gabor: Gabor kernel

*** kernel** Custom kernel.

Option	Description
Kernel.SizeX [0]	Width of the kernel (number of columns)
Kernel.SizeY [0]	Height of the kernel (number of rows)
Kernel.Mat	List of row-major ordered coefficients of the kernel

If both `Kernel.SizeX` and `Kernel.SizeY` are 0, the kernel is assumed to be square.

Gaussian kernel Gaussian kernel.

Option [default value]	Description
Kernel.SizeX	Width of the kernel (number of columns)
Kernel.SizeY	Height of the kernel (number of rows)
Kernel.Positive [1]	If true, the center of the kernel is positive
Kernel.Sigma [$\sqrt{2.0}$]	Sigma of the kernel

LoG kernel Laplacian Of Gaussian kernel.

Option [default value]	Description
Kernel.SizeX	Width of the kernel (number of columns)
Kernel.SizeY	Height of the kernel (number of rows)
Kernel.Positive [1]	If true, the center of the kernel is positive
Kernel.Sigma [$\sqrt{2.0}$]	Sigma of the kernel

DoG kernel Difference Of Gaussian kernel kernel.

Option [default value]	Description
Kernel.SizeX	Width of the kernel (number of columns)
Kernel.SizeY	Height of the kernel (number of rows)
Kernel.Positive [1]	If true, the center of the kernel is positive
Kernel.Sigma1 [2.0]	Sigma1 of the kernel
Kernel.Sigma2 [1.0]	Sigma2 of the kernel

Gabor kernel Gabor kernel.

Option [default value]	Description
Kernel.SizeX	Width of the kernel (number of columns)
Kernel.SizeY	Height of the kernel (number of rows)
Kernel.Theta	Theta of the kernel
Kernel.Sigma [$\sqrt{2.0}$]	Sigma of the kernel
Kernel.Lambda [10.0]	Lambda of the kernel
Kernel.Psi [$\pi/2.0$]	Psi of the kernel
Kernel.Gamma [0.5]	Gamma of the kernel

FlipTransformation Image flip transformation.

Option [default value]	Description
HorizontalFlip [0]	If true, flip the image horizontally
VerticalFlip [0]	If true, flip the image vertically
RandomHorizontalFlip [0]	If true, randomly flip the image horizontally
RandomVerticalFlip [0]	If true, randomly flip the image vertically

N2D2 IP only

GradientFilterTransformation Compute image gradient.

Option [default value]	Description
Scale [1.0]	Scale to apply to the computed gradient
Delta [0.0]	Bias to add to the computed gradient
GradientFilter [Sobel]	Filter type to use for computing the gradient. Possible options are: Sobel, Scharr and Laplacian
KernelSize [3]	Size of the filter kernel (has no effect when using the Scharr filter, which kernel size is always 3x3)
ApplyToLabels [0]	If true, use the computed gradient to filter the image label and ignore pixel areas where the gradient is below the Threshold . In this case, only the labels are modified, not the image
InvThreshold [0]	If true, ignored label pixels will be the ones with a low gradient (low contrasted areas)
Threshold [0.5]	Threshold applied on the image gradient
Label []	List of labels to filter (space-separated)
GradientScale [1.0]	Rescale the image by this factor before applying the gradient and the threshold, then scale it back to filter the labels

N2D2 IP only

LabelSliceExtractionTransformation Extract a slice from an image belonging to a given label.

Option [default value]	Description
Width	Width of the slice to extract
Height	Height of the slice to extract
Label [-1]	Slice should belong to this label ID. If -1, the label ID is random
RandomRotation [0]	If true, extract randomly rotated slices
RandomRotationRange [0.0 360.0]	Range of the random rotations, in degrees, counterclockwise (if RandomRotation is enabled)
SlicesMargin [0]	Positive or negative, indicates the margin around objects that can be extracted in the slice
KeepComposite [0]	If false, the 2D label image is reduced to a single value corresponding to the extracted object label (useful for patches classification tasks). Note that if SlicesMargin is > 0 , the 2D label image may contain other labels before reduction. For pixel-wise segmentation tasks, set KeepComposite to true.

This transformation is useful to learn sparse object occurrences in a lot of background. If the dataset is very unbalanced towards background, this transformation will ensure that the learning is done on a more balanced set of every labels, regardless of their actual pixel-wise ratio.

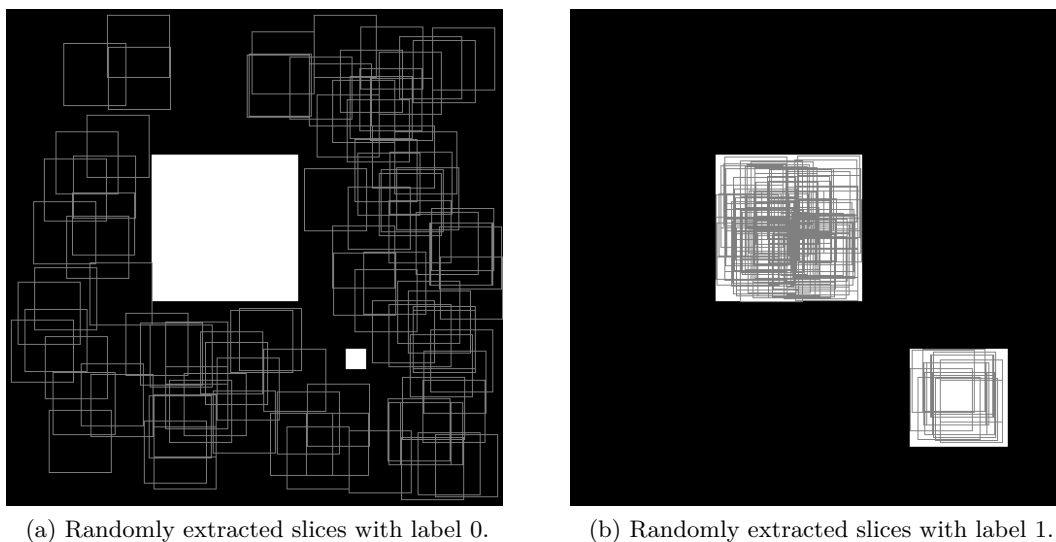


Figure 12: Illustration of the working behavior of **LabelSliceExtractionTransformation** with **SlicesMargin** = 0.

When **SlicesMargin** is 0, only slices that fully include a given label are extracted, as shown in figure 12. The behavior with **SlicesMargin** < 0 is illustrated in figure 13. Note that setting a negative **SlicesMargin** larger in absolute value than **Width**/2 or **Height**/2 will lead in some (random) cases in incorrect slice labels in respect to the majority pixel label in the slice.

MagnitudePhaseTransformation Compute the magnitude and phase of a complex two channels input data, with the first channel x being the real part and the second channel y the imaginary part. The resulting data is two channels, the first one with the magnitude and the second one with the phase.

Option [default value]	Description
LogScale [0]	If true, compute the magnitude in log scale



Figure 13: Illustration of the working behavior of **LabelSliceExtractionTransformation** with **SlicesMargin** = -32.

The magnitude is:

$$M_{i,j} = \sqrt{x_{i,j}^2 + x_{i,j}^2}$$

If **LogScale** = 1, compute $M'_{i,j} = \log(1 + M_{i,j})$.

The phase is:

$$\theta_{i,j} = \text{atan2}(y_{i,j}, x_{i,j})$$

N2D2 IP only MorphologicalReconstructionTransformation Apply a morphological reconstruction transformation to the image. This transformation is also useful for post-processing.

Option [default value]	Description
Operation	Morphological operation to apply. Can be: ReconstructionByErosion : reconstruction by erosion operation ReconstructionByDilation : reconstruction by dilation operation OpeningByReconstruction : opening by reconstruction operation ClosingByReconstruction : closing by reconstruction operation
Size	Size of the structuring element
ApplyToLabels [0]	If true, apply the transformation to the labels instead of the image
Shape [Rectangular]	Shape of the structuring element used for morphology operations. Can be Rectangular , Elliptic or Cross .
NbIterations [1]	Number of times erosion and dilation are applied for opening and closing reconstructions

N2D2 IP only MorphologyTransformation Apply a morphology transformation to the image. This transformation is also useful for post-processing.

Option [default value]	Description
Operation	Morphological operation to apply. Can be: Erode : erode operation ($= \text{erode}(\text{src})$) Dilate : dilate operation ($= \text{dilate}(\text{src})$) Opening : opening operation ($\text{open}(\text{src}) = \text{dilate}(\text{erode}(\text{src}))$) Closing : closing operation ($\text{close}(\text{src}) = \text{erode}(\text{dilate}(\text{src}))$) Gradient : morphological gradient ($= \text{dilate}(\text{src}) - \text{erode}(\text{src})$) TopHat : top hat ($= \text{src} - \text{open}(\text{src})$) BlackHat : black hat ($= \text{close}(\text{src}) - \text{src}$)
Size	Size of the structuring element
ApplyToLabels [0]	If true, apply the transformation to the labels instead of the image
Shape [Rectangular]	Shape of the structuring element used for morphology operations. Can be Rectangular , Elliptic or Cross .
NbIterations [1]	Number of times erosion and dilation are applied

NormalizeTransformation Normalize the image.

Option [default value]	Description
Norm [MinMax]	Norm type, can be: L1 : L1 normalization L2 : L2 normalization Linf : Linf normalization MinMax : min-max normalization
NormValue [1.0]	Norm value (for L1, L2 and Linf) Such that $\ data\ _{L_p} = \text{NormValue}$
NormMin [0.0]	Min value (for MinMax only) Such that $\min(data) = \text{NormMin}$
NormMax [1.0]	Max value (for MinMax only) Such that $\max(data) = \text{NormMax}$
PerChannel [0]	If true, normalize each channel individually

PadCropTransformation Pad/crop the image to a specified size.

Option [default value]	Description
Width	Width of the padded/cropped image
Height	Height of the padded/cropped image
BorderType [MinusOneReflectBorder]	Border type used when padding. Possible values: ConstantBorder : pad with BorderValue ReplicateBorder : last element is replicated throughout, like aaaaaa abcdefgh hhhhhhh ReflectBorder : border will be mirror reflection of the border elements, like fedcba abcdefgh hgfedcb WrapBorder : it will look like cdefgh abcdefgh abcdefg MinusOneReflectBorder : same as ReflectBorder but with a slight change, like gfedcb abcdefgh gfedcba MeanBorder : pad with the mean color of the image
BorderValue [0.0 0.0 0.0]	Background color used when padding with BorderType is ConstantBorder

N2D2 IP only RandomAffineTransformation Apply a global random affine transformation to the values of the image.

Option [default value]	Description
GainRange [1.0 1.0]	Random gain (α) range (identical for all channels)
GainRange[*] [1.0 1.0]	Random gain (α) range for channel *. Mutually exclusive with GainRange . If any specified, a different random gain will always be sampled for each channel. Default gain is 1.0 (no gain) for missing channels
	The gain control the <i>contrast</i> of the image
BiasRange [0.0 0.0]	Random bias (β) range (identical for all channels)
BiasRange[*] [0.0 0.0]	Random bias (β) range for channel *. Mutually exclusive with BiasRange . If any specified, a different random bias will always be sampled for each channel. Default bias is 0.0 (no bias) for missing channels
	The bias control the <i>brightness</i> of the image
GammaRange [1.0 1.0]	Random gamma (γ) range (identical for all channels)
GammaRange[*] [1.0 1.0]	Random gamma (γ) range for channel *. Mutually exclusive with GammaRange . If any specified, a different random gamma will always be sampled for each channel. Default gamma is 1.0 (no change) for missing channels
	The gamma control more or less the <i>exposure</i> of the image
GainVarProb [1.0]	Probability to have a gain variation for each channel. If only one value is specified, the same probability applies to all the channels. In this case, the same gain variation will be sampled for all the channels only if a single range is specified for all the channels using GainRange . If more than one value is specified, a different random gain will always be sampled for each channel, even if the probabilities and ranges are identical
BiasVarProb [1.0]	Probability to have a bias variation for each channel. If only one value is specified, the same probability applies to all the channels. In this case, the same bias variation will be sampled for all the channels only if a single range is specified for all the channels using BiasRange . If more than one value is specified, a different random bias will always be sampled for each channel, even if the probabilities and ranges are identical
GammaVarProb [1.0]	Probability to have a gamma variation for each channel. If only one value is specified, the same probability applies to all the channels. In this case, the same gamma variation will be sampled for all the channels only if a single range is specified for all the channels using GammaRange . If more than one value is specified, a different random gamma will always be sampled for each channel, even if the probabilities and ranges are identical
DisjointGamma [0]	If true, gamma variation and gain/bias variation are mutually exclusive. The probability to have a random gamma variation is therefore GammaVarProb and the probability to have a gain/bias variation is $1 - \text{GammaVarProb}$.
ChannelsMask []	If not empty, specifies on which channels the transformation is applied. For example, to apply the transformation only to the first and third channel, set ChannelsMask to 1 0 1

The equation of the transformation is:

$$S = \begin{cases} \text{numeric_limits}<T>::\text{max}() & \text{if is_integer}<T> \\ 1.0 & \text{otherwise} \end{cases}$$

$$v(i, j) = \text{cv}::\text{saturate_cast}<T> \left(\alpha \left(\frac{v(i, j)}{S} \right)^\gamma S + \beta.S \right)$$

RangeAffineTransformation Apply an affine transformation to the values of the image.

Option [default value]	Description
FirstOperator	First operator, can be Plus, Minus, Multiplies, Divides
FirstValue	
SecondOperator [Plus]	Second operator, can be Plus, Minus, Multiplies, Divides
SecondValue [0.0]	

The final operation is the following:

$$f(x) = (x \stackrel{o}{op}_{1st} val_{1st}) \stackrel{o}{op}_{2nd} val_{2nd}$$

N2D2 IP only **RangeClippingTransformation** Clip the value range of the image.

Option [default value]	Description
RangeMin [$\min(data)$]	Image values below RangeMin are clipped to 0
RangeMax [$\max(data)$]	

RescaleTransformation Rescale the image to a specified size.

Option [default value]	Description
Width	Width of the rescaled image
Height	
KeepAspectRatio [0]	If true, keeps the aspect ratio of the image
ResizeToFit [1]	

ReshapeTransformation Reshape the data to a specified size.

Option [default value]	Description
NbRows	New number of rows
NbCols [0]	
NbChannels [0]	

N2D2 IP only **SliceExtractionTransformation** Extract a slice from an image.

Option [default value]	Description
Width	Width of the slice to extract
Height	Height of the slice to extract
OffsetX [0]	X offset of the slice to extract
OffsetY [0]	Y offset of the slice to extract
RandomOffsetX [0]	If true, the X offset is chosen randomly
RandomOffsetY [0]	If true, the Y offset is chosen randomly
RandomRotation [0]	If true, extract randomly rotated slices
RandomRotationRange [0.0 360.0]	Range of the random rotations, in degrees, counterclockwise (if RandomRotation is enabled)
RandomScaling [0]	If true, extract randomly scaled slices
RandomScalingRange [0.8 1.2]	Range of the random scaling (if RandomRotation is enabled)
AllowPadding [0]	If true, zero-padding is allowed if the image is smaller than the slice to extract
BorderType [MinusOneReflectBorder]	Border type used when padding. Possible values: ConstantBorder : pad with BorderValue ReplicateBorder : last element is replicated throughout, like aaaaaa abcdefgh hhhhhhh ReflectBorder : border will be mirror reflection of the border elements, like fedcba abcdefgh hgfedcb WrapBorder : it will look like cdefgh abcdefgh abcdefg MinusOneReflectBorder : same as ReflectBorder but with a slight change, like gfedcb abcdefgh gfedcba MeanBorder : pad with the mean color of the image
BorderValue [0.0 0.0 0.0]	Background color used when padding with BorderType is ConstantBorder

ThresholdTransformation Apply a thresholding transformation to the image. This transformation is also useful for post-processing.

Option [default value]	Description
Threshold	Threshold value
OtsuMethod [0]	Use Otsu's method to determine the optimal threshold (if true, the Threshold value is ignored)
Operation [Binary]	Thresholding operation to apply. Can be: Binary BinaryInverted Truncate ToZero ToZeroInverted
MaxValue [1.0]	Max. value to use with Binary and BinaryInverted operations

TrimTransformation Trim the image.

Option [default value]	Description
NbLevels	Number of levels for the color discretization of the image
Method [Discretize]	Possible values are: Reduce : discretization using K-means Discretize : simple discretization

WallisFilterTransformation Apply Wallis filter to the image.

Option [default value]	Description
Size	Size of the filter
Mean [0.0]	Target mean value
StdDev [1.0]	Target standard deviation
PerChannel [0]	If true, apply Wallis filter to each channel individually (this parameter is meaningful only if Size is 0)

4.7 Network layers

4.7.1 Layer definition

Common set of parameters for any kind of layer.

Option [default value]	Description
Input	Name of the section(s) for the input layer(s). Comma separated
Type	Type of the layer. Can be any of the type described below
Model [DefaultModel]	Layer model to use
DataType [DefaultDataType]	Layer data type to use. Please note that some layers may not support every data type.
ConfigSection []	Name of the configuration section for layer

To specify that the back-propagated error must be computed at the output of a given layer (generally the last layer, or output layer), one must add a target section named *LayerName.Target*:

```
...
[LayerName.Target]
TargetValue=1.0 ; default: 1.0
DefaultValue=0.0 ; default: -1.0
```

4.7.2 Weight fillers

Fillers to initialize weights and biases in the different type of layer.

Usage example:

```
[conv1]
...
WeightsFiller=NormalFiller
WeightsFiller.Mean=0.0
WeightsFiller.StdDev=0.05
...
```

The initial weights distribution for each layer can be checked in the *weights_init* folder, with an example shown in figure 14.

ConstantFiller Fill with a constant value.

Option	Description
<i>FillerName.Value</i>	Value for the filling




Figure 14: Initial weights distribution of a layer using a normal distribution (**NormalFiller**) with a 0 mean and a 0.05 standard deviation.

HeFiller Fill with an normal distribution with normalized variance taking into account the rectifier nonlinearity (He et al., 2015). This filler is sometimes referred as MSRA filler.

Option [default value]	Description
<i>FillerName.VarianceNorm</i> [FanIn]	Normalization, can be FanIn, Average or FanOut
<i>FillerName.Scaling</i> [1.0]	Scaling factor

Use a normal distribution with standard deviation $\sqrt{\frac{2.0}{n}}$.

- $n = \text{fan-in}$ with FanIn, resulting in $\text{Var}(W) = \frac{2}{\text{fan-in}}$
- $n = \frac{(\text{fan-in} + \text{fan-out})}{2}$ with Average, resulting in $\text{Var}(W) = \frac{4}{\text{fan-in} + \text{fan-out}}$
- $n = \text{fan-out}$ with FanOut, resulting in $\text{Var}(W) = \frac{2}{\text{fan-out}}$

NormalFiller Fill with a normal distribution.

Option [default value]	Description
<i>FillerName.Mean</i> [0.0]	Mean value of the distribution
<i>FillerName.StdDev</i> [1.0]	Standard deviation of the distribution

UniformFiller Fill with an uniform distribution.

Option [default value]	Description
<i>FillerName.Min</i> [0.0]	Min. value
<i>FillerName.Max</i> [1.0]	Max. value

XavierFiller Fill with an uniform distribution with normalized variance (Glorot and Bengio, 2010).

Option [default value]	Description
<i>FillerName.VarianceNorm</i> [FanIn]	Normalization, can be FanIn, Average or FanOut
<i>FillerName.Distribution</i> [Uniform]	Distribution, can be Uniform or Normal
<i>FillerName.Scaling</i> [1.0]	Scaling factor

Use an uniform distribution with interval $[-scale, scale]$, with $scale = \sqrt{\frac{3.0}{n}}$.

- $n = fan-in$ with FanIn, resulting in $Var(W) = \frac{1}{fan-in}$
- $n = \frac{(fan-in+fan-out)}{2}$ with Average, resulting in $Var(W) = \frac{2}{fan-in+fan-out}$
- $n = fan-out$ with FanOut, resulting in $Var(W) = \frac{1}{fan-out}$

4.7.3 Weight solvers

SGDSolver_Frame SGD Solver for Frame models.

Option [default value]	Description
<i>SolverName.LearningRate</i> [0.01]	Learning rate
<i>SolverName.Momentum</i> [0.0]	Momentum
<i>SolverName.Decay</i> [0.0]	Decay
<i>SolverName.LearningRatePolicy</i> [None]	Learning rate decay policy. Can be any of None, StepDecay, ExponentialDecay, InvTDecay, PolyDecay
<i>SolverName.LearningRateStepSize</i> [1]	Learning rate step size (in number of stimuli)
<i>SolverName.LearningRateDecay</i> [0.1]	Learning rate decay
<i>SolverName.Clamping</i> [0]	If true, clamp the weights and bias between -1 and 1
<i>SolverName.Power</i> [0.0]	Polynomial learning rule power parameter
<i>SolverName.MaxIterations</i> [0.0]	Polynomial learning rule maximum number of iterations

The learning rate decay policies are the following:

- **StepDecay:** every *SolverName.LearningRateStepSize* stimuli, the learning rate is reduced by a factor *SolverName.LearningRateDecay*;
- **ExponentialDecay:** the learning rate is $\alpha = \alpha_0 \exp(-kt)$, with α_0 the initial learning rate *SolverName.LearningRate*, k the rate decay *SolverName.LearningRateDecay* and t the step number (one step every *SolverName.LearningRateStepSize* stimuli);
- **InvTDecay:** the learning rate is $\alpha = \alpha_0 / (1 + kt)$, with α_0 the initial learning rate *SolverName.LearningRate*, k the rate decay *SolverName.LearningRateDecay* and t the step number (one step every *SolverName.LearningRateStepSize* stimuli).

- **InvDecay**: the learning rate is $\alpha = \alpha_0 * (1 + kt)^{-n}$, with α_0 the initial learning rate *SolverName.LearningRate*, k the rate decay *SolverName.LearningRateDecay*, t the current iteration and n the power parameter *SolverName.Power*
- **PolyDecay**: the learning rate is $\alpha = \alpha_0 * (1 - \frac{k}{t})^n$, with α_0 the initial learning rate *SolverName.LearningRate*, k the current iteration, t the maximum number of iteration *SolverName.MaxIterations* and n the power parameter *SolverName.Power*

SGDSolver_Frame_CUDA SGD Solver for **Frame_CUDA** models.

Option [default value]	Description
<i>SolverName.LearningRate</i> [0.01]	Learning rate
<i>SolverName.Momentum</i> [0.0]	Momentum
<i>SolverName.Decay</i> [0.0]	Decay
<i>SolverName.LearningRatePolicy</i> [None]	Learning rate decay policy. Can be any of None , StepDecay , ExponentialDecay , InvTDecay
<i>SolverName.LearningRateStepSize</i> [1]	Learning rate step size (in number of stimuli)
<i>SolverName.LearningRateDecay</i> [0.1]	Learning rate decay
<i>SolverName.Clamping</i> [0]	If true, clamp the weights and bias between -1 and 1

The learning rate decay policies are identical to the ones in the **SGDSolver_Frame** solver.

AdamSolver_Frame Adam Solver for **Frame** models ([Kingma and Ba, 2014](#)).

Option [default value]	Description
<i>SolverName.LearningRate</i> [0.001]	Learning rate (stepsize)
<i>SolverName.Beta1</i> [0.9]	Exponential decay rate of these moving average of the first moment
<i>SolverName.Beta2</i> [0.999]	Exponential decay rate of these moving average of the second moment
<i>SolverName.Epsilon</i> [1.0e-8]	Epsilon

AdamSolver_Frame_CUDA Adam Solver for **Frame_CUDA** models ([Kingma and Ba, 2014](#)).

Option [default value]	Description
<i>SolverName.LearningRate</i> [0.001]	Learning rate (stepsize)
<i>SolverName.Beta1</i> [0.9]	Exponential decay rate of these moving average of the first moment
<i>SolverName.Beta2</i> [0.999]	Exponential decay rate of these moving average of the second moment
<i>SolverName.Epsilon</i> [1.0e-8]	Epsilon

4.7.4 Activation functions

Activation function to be used at the output of layers.

Usage example:

```
[conv1]
...
ActivationFunction=Rectifier
ActivationFunction.LeakSlope=0.01
ActivationFunction.Clipping=20
...
```

Logistic Logistic activation function.

LogisticWithLoss Logistic with loss activation function.

Rectifier Rectifier or ReLU activation function.

Option [default value]	Description
ActivationFunction.LeakSlope [0.0]	Leak slope for negative inputs
ActivationFunction.Clipping [0.0]	Clipping value for positive outputs

Saturation Saturation activation function.

Softplus Softplus activation function.

Tanh Tanh activation function.
Computes $y = \tanh(\alpha x)$.

Option [default value]	Description
ActivationFunction.Alpha [1.0]	α parameter

TanhLeCun Tanh activation function with an α parameter of $1.7159 \times (2.0/3.0)$.

4.7.5 Anchor

Anchor layer for Faster R-CNN or Single Shot Detector.

Option [default value]	Description
Input	This layer takes one or two inputs. The total number of input channels must be <code>ScoresCls + 4</code> , with <code>ScoresCls</code> being equal to 1 or 2.
Anchor[*]	Anchors definition. For each anchor, there must be two space-separated values: the root area and the aspect ratio.
ScoresCls	Number of classes per anchor. Must be 1 (if the scores input uses logistic regression) or 2 (if the scores input is a two-class softmax layer)
FeatureMapWidth [StimuliProvider.Width]	Reference width use to scale anchors coordinate.
FeatureMapHeight [StimuliProvider.Height]	Reference height use to scale anchors coordinate.

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
PositiveIoU [0.7]	<i>all Frame</i>	Assign a positive label for anchors whose IoU overlap is higher than <code>PositiveIoU</code> with any ground-truth box
NegativeIoU [0.3]	<i>all Frame</i>	Assign a negative label for non-positive anchors whose IoU overlap is lower than <code>NegativeIoU</code> for all ground-truth boxes
LossLambda [10.0]	<i>all Frame</i>	Balancing parameter λ
LossPositiveSample [128]	<i>all Frame</i>	Number of random positive samples for the loss computation
LossNegativeSample [128]	<i>all Frame</i>	Number of random negative samples for the loss computation

Usage example:

```
; RPN network: cls layer
[scores]
Input=...
Type=Conv
KernelWidth=1
KernelHeight=1
; 18 channels for 9 anchors
NbOutputs=18
...

[scores.softmax]
Input=scores
Type=Softmax
NbOutputs=[scores]NbOutputs
WithLoss=1

; RPN network: coordinates layer
[coordinates]
Input=...
Type=Conv
KernelWidth=1
KernelHeight=1
; 36 channels for 4 coordinates x 9 anchors
NbOutputs=36
...

; RPN network: anchors
[anchors]
Input=scores.softmax,coordinates
Type=Anchor
ScoresCls=2 ; using a two-class softmax for the scores
Anchor[0]=32 1.0
Anchor[1]=48 1.0
Anchor[2]=64 1.0
Anchor[3]=80 1.0
Anchor[4]=96 1.0
Anchor[5]=112 1.0
Anchor[6]=128 1.0
Anchor[7]=144 1.0
Anchor[8]=160 1.0
ConfigSection=anchors.config
```

```
[anchors.config]
PositiveIoU=0.7
NegativeIoU=0.3
LossLambda=1.0
```

Outputs remapping Outputs remapping allows to convert *scores* and *coordinates* output feature maps layout from another ordering that the one used in the N2D2 `Anchor` layer, during weights import/export.

For example, lets consider that the imported weights corresponds to the following output feature maps ordering:

```
0 anchor[0].y
1 anchor[0].x
2 anchor[0].h
3 anchor[0].w
4 anchor[1].y
5 anchor[1].x
6 anchor[1].h
7 anchor[1].w
8 anchor[2].y
9 anchor[2].x
10 anchor[2].h
11 anchor[2].w
```

The output feature maps ordering required by the `Anchor` layer is:

```
0 anchor[0].x
1 anchor[1].x
2 anchor[2].x
3 anchor[0].y
4 anchor[1].y
5 anchor[2].y
6 anchor[0].w
7 anchor[1].w
8 anchor[2].w
9 anchor[0].h
10 anchor[1].h
11 anchor[2].h
```

The feature maps ordering can be changed during weights import/export:

```
; RPN network: coordinates layer
[coordinates]
Input=...
Type=Conv
KernelWidth=1
KernelHeight=1
; 36 channels for 4 coordinates x 9 anchors
NbOutputs=36
...
ConfigSection=coordinates.config

[coordinates.config]
WeightsExportFormat=HWCO ; Weights format used by TensorFlow
OutputsRemap=1:4,0:4,3:4,2:4
```

4.7.6 Conv

Convolutional layer.

Option [default value]	Description
------------------------	-------------

KernelWidth	Width of the kernels
KernelHeight	Height of the kernels
KernelDepth []	Depth of the kernels (implies 3D kernels)
OR	
KernelSize []	Kernels size (implies 2D square kernels)
OR	
KernelDims []	List of space-separated dimensions for N-D kernels
NbOutputs	Number of output channels
SubSampleX [1]	X-axis subsampling factor of the output feature maps
SubSampleY [1]	Y-axis subsampling factor of the output feature maps
SubSampleZ []	Z-axis subsampling factor of the output feature maps
OR	
SubSample [1]	Subsampling factor of the output feature maps
OR	
SubSampleDims []	List of space-separated subsampling dimensions for N-D kernels
StrideX [1]	X-axis stride of the kernels
StrideY [1]	Y-axis stride of the kernels
StrideZ []	Z-axis stride of the kernels
OR	
Stride [1]	Stride of the kernels
OR	
StrideDims []	List of space-separated stride dimensions for N-D kernels
PaddingX [0]	X-axis input padding
PaddingY [0]	Y-axis input padding
PaddingZ []	Z-axis input padding
OR	
Padding [0]	Input padding
OR	
PaddingDims []	List of space-separated padding dimensions for N-D kernels
DilationX [1]	X-axis dilation of the kernels
DilationY [1]	Y-axis dilation of the kernels
DilationZ []	Z-axis dilation of the kernels
OR	
Dilation [1]	Dilation of the kernels
OR	
DilationDims []	List of space-separated dilation dimensions for N-D kernels
ActivationFunction [Tanh]	Activation function. Can be any of Logistic , LogisticWithLoss , Rectifier , Softplus , TanhLeCun , Linear , Saturation OR Tanh
WeightsFiller [NormalFiller(0.0, 0.05)]	Weights initial values filler
BiasFiller [NormalFiller(0.0, 0.05)]	Biases initial values filler
Mapping.NbGroups []	Mapping: number of groups (mutually exclusive with all other Mapping.* options)
Mapping.ChannelsPerGroup []	Mapping: number of channels per group (mutually exclusive with all other Mapping.* options)
Mapping.SizeX [1]	Mapping canvas pattern default width
Mapping.SizeY [1]	Mapping canvas pattern default height
Mapping.Size [1]	Mapping canvas pattern default size (mutually exclusive with Mapping.SizeX and Mapping.SizeY)

<code>Mapping.StrideX</code> [1]	Mapping canvas default X-axis step
<code>Mapping.StrideY</code> [1]	Mapping canvas default Y-axis step
<code>Mapping.Stride</code> [1]	Mapping canvas default step (mutually exclusive with <code>Mapping.StrideX</code> and <code>Mapping.StrideY</code>)
<code>Mapping.OffsetX</code> [0]	Mapping canvas default X-axis offset
<code>Mapping.OffsetY</code> [0]	Mapping canvas default Y-axis offset
<code>Mapping.Offset</code> [0]	Mapping canvas default offset (mutually exclusive with <code>Mapping.OffsetX</code> and <code>Mapping.OffsetY</code>)
<code>Mapping.NbIterations</code> [0]	Mapping canvas pattern default number of iterations (0 means no limit)
<code>Mapping(in).SizeX</code> [1]	Mapping canvas pattern default width for input layer <code>in</code>
<code>Mapping(in).SizeY</code> [1]	Mapping canvas pattern default height for input layer <code>in</code>
<code>Mapping(in).Size</code> [1]	Mapping canvas pattern default size for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).SizeX</code> and <code>Mapping(in).SizeY</code>)
<code>Mapping(in).StrideX</code> [1]	Mapping canvas default X-axis step for input layer <code>in</code>
<code>Mapping(in).StrideY</code> [1]	Mapping canvas default Y-axis step for input layer <code>in</code>
<code>Mapping(in).Stride</code> [1]	Mapping canvas default step for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).StrideX</code> and <code>Mapping(in).StrideY</code>)
<code>Mapping(in).OffsetX</code> [0]	Mapping canvas default X-axis offset for input layer <code>in</code>
<code>Mapping(in).OffsetY</code> [0]	Mapping canvas default Y-axis offset for input layer <code>in</code>
<code>Mapping(in).Offset</code> [0]	Mapping canvas default offset for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).OffsetX</code> and <code>Mapping(in).OffsetY</code>)
<code>Mapping(in).NbIterations</code> [0]	Mapping canvas pattern default number of iterations for input layer <code>in</code> (0 means no limit)
<code>WeightsSharing</code> []	Share the weights with an other layer
<code>BiasesSharing</code> []	Share the biases with an other layer

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
<code>NoBias</code> [0]	<i>all Frame</i>	If true, don't use bias
<code>Solvers.*</code>	<i>all Frame</i>	Any solver parameters
<code>WeightsSolver.*</code>	<i>all Frame</i>	Weights solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>BiasSolver.*</code>	<i>all Frame</i>	Bias solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>WeightsExportFormat</code> [OCHW]	<i>all Frame</i>	Weights import/export format. Can be <code>oCHW</code> or <code>OCHW</code> , with <code>o</code> the output feature map, <code>c</code> the input feature map (channel), <code>H</code> the kernel row and <code>w</code> the kernel column, in the order of the outermost dimension (in the leftmost position) to the innermost dimension (in the rightmost position)
<code>WeightsExportFlip</code> [0]	<i>all Frame</i>	If true, import/export flipped kernels

Configuration parameters (*Spike models*)

Experimental option (implementation may be wrong or susceptible to change)

Option [default value]	Model(s)	Description
IncomingDelay [1 TimePs ;100 TimeFs]	<i>all Spike</i>	Synaptic incoming delay w_{delay}
Threshold [1.0]	Spike, Spike_RRAM	Threshold of the neuron I_{thres}
BipolarThreshold [1]	Spike, Spike_RRAM	If true, the threshold is also applied to the absolute value of negative values (generating negative spikes)
Leak [0.0]	Spike, Spike_RRAM	Neural leak time constant τ_{leak} (if 0, no leak)
Refractory [0.0]	Spike, Spike_RRAM	Neural refractory period T_{refrac}
WeightsRelInit [0.0;0.05]	Spike	Relative initial synaptic weight w_{init}
WeightsMinMean [1;0.1]	Spike_RRAM	Mean minimum synaptic weight w_{min}
WeightsMaxMean [100;10.0]	Spike_RRAM	Mean maximum synaptic weight w_{max}
WeightsMinVarSlope [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMinVarOrigin [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMaxVarSlope [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMaxVarOrigin [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsSetProba [1.0]	Spike_RRAM	Intrinsic SET switching probability P_{SET} (upon receiving a SET programming pulse). Assuming uniform statistical distribution (not well supported by experiments on RRAM)
WeightsResetProba [1.0]	Spike_RRAM	Intrinsic RESET switching probability P_{RESET} (upon receiving a RESET programming pulse). Assuming uniform statistical distribution (not well supported by experiments on RRAM)
SynapticRedundancy [1]	Spike_RRAM	Synaptic redundancy (number of RRAM device per synapse)
BipolarWeights [0]	Spike_RRAM	Bipolar weights
BipolarIntegration [0]	Spike_RRAM	Bipolar integration
LtpProba [0.2]	Spike_RRAM	Extrinsic STDP LTP probability (cumulative with intrinsic SET switching probability P_{SET})
LtdProba [0.1]	Spike_RRAM	Extrinsic STDP LTD probability (cumulative with intrinsic RESET switching probability P_{RESET})
Stdpltp [1000 TimePs]	Spike_RRAM	STDP LTP time window T_{LTP}
InhibitRefractory [0 TimePs]	Spike_RRAM	Neural lateral inhibition period $T_{inhibit}$
EnableStdpl [1]	Spike_RRAM	If false, STDP is disabled (no synaptic weight change)
RefractoryIntegration [1]	Spike_RRAM	If true, reset the integration to 0 during the refractory period
DigitalIntegration [0]	Spike_RRAM	If false, the analog value of the devices is integrated, instead of their binary value

4.7.7 Deconv

Deconvolutionlayer.

Option [default value]	Description
------------------------	-------------

KernelWidth	Width of the kernels
KernelHeight	Height of the kernels
KernelDepth []	Depth of the kernels (implies 3D kernels)
OR	
KernelSize []	Kernels size (implies 2D square kernels)
OR	
KernelDims []	List of space-separated dimensions for N-D kernels
NbOutputs	Number of output channels
StrideX [1]	X-axis stride of the kernels
StrideY [1]	Y-axis stride of the kernels
StrideZ []	Z-axis stride of the kernels
OR	
Stride [1]	Stride of the kernels
OR	
StrideDims []	List of space-separated stride dimensions for N-D kernels
PaddingX [0]	X-axis input padding
PaddingY [0]	Y-axis input padding
PaddingZ []	Z-axis input padding
OR	
Padding [0]	Input padding
OR	
PaddingDims []	List of space-separated padding dimensions for N-D kernels
DilationX [1]	X-axis dilation of the kernels
DilationY [1]	Y-axis dilation of the kernels
DilationZ []	Z-axis dilation of the kernels
OR	
Dilation [1]	Dilation of the kernels
OR	
DilationDims []	List of space-separated dilation dimensions for N-D kernels
ActivationFunction [Tanh]	Activation function. Can be any of Logistic, LogisticWithLoss, Rectifier, Softplus, TanhLeCun, Linear, Saturation OR Tanh
WeightsFiller [NormalFiller(0.0, 0.05)]	Weights initial values filler
BiasFiller [NormalFiller(0.0, 0.05)]	Biases initial values filler
Mapping.NbGroups []	Mapping: number of groups (mutually exclusive with all other Mapping.* options)
Mapping.ChannelsPerGroup []	Mapping: number of channels per group (mutually exclusive with all other Mapping.* options)
Mapping.SizeX [1]	Mapping canvas pattern default width
Mapping.SizeY [1]	Mapping canvas pattern default height
Mapping.Size [1]	Mapping canvas pattern default size (mutually exclusive with Mapping.SizeX and Mapping.SizeY)
Mapping.StrideX [1]	Mapping canvas default X-axis step
Mapping.StrideY [1]	Mapping canvas default Y-axis step
Mapping.Stride [1]	Mapping canvas default step (mutually exclusive with Mapping.StrideX and Mapping.StrideY)
Mapping.OffsetX [0]	Mapping canvas default X-axis offset
Mapping.OffsetY [0]	Mapping canvas default Y-axis offset
Mapping.Offset [0]	Mapping canvas default offset (mutually exclusive with Mapping.OffsetX and Mapping.OffsetY)

<code>Mapping.NbIterations</code> [0]	Mapping canvas pattern default number of iterations (0 means no limit)
<code>Mapping(in).SizeX</code> [1]	Mapping canvas pattern default width for input layer <code>in</code>
<code>Mapping(in).SizeY</code> [1]	Mapping canvas pattern default height for input layer <code>in</code>
<code>Mapping(in).Size</code> [1]	Mapping canvas pattern default size for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).SizeX</code> and <code>Mapping(in).SizeY</code>)
<code>Mapping(in).StrideX</code> [1]	Mapping canvas default X-axis step for input layer <code>in</code>
<code>Mapping(in).StrideY</code> [1]	Mapping canvas default Y-axis step for input layer <code>in</code>
<code>Mapping(in).Stride</code> [1]	Mapping canvas default step for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).StrideX</code> and <code>Mapping(in).StrideY</code>)
<code>Mapping(in).OffsetX</code> [0]	Mapping canvas default X-axis offset for input layer <code>in</code>
<code>Mapping(in).OffsetY</code> [0]	Mapping canvas default Y-axis offset for input layer <code>in</code>
<code>Mapping(in).Offset</code> [0]	Mapping canvas default offset for input layer <code>in</code> (mutually exclusive with <code>Mapping(in).OffsetX</code> and <code>Mapping(in).OffsetY</code>)
<code>Mapping(in).NbIterations</code> [0]	Mapping canvas pattern default number of iterations for input layer <code>in</code> (0 means no limit)
<code>WeightsSharing</code> []	Share the weights with an other layer
<code>BiasesSharing</code> []	Share the biases with an other layer

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
<code>NoBias</code> [0]	<i>all Frame</i>	If true, don't use bias
<code>BackPropagate</code> [1]	<i>all Frame</i>	If true, enable backpropagation
<code>Solvers.*</code>	<i>all Frame</i>	Any solver parameters
<code>WeightsSolver.*</code>	<i>all Frame</i>	Weights solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>BiasSolver.*</code>	<i>all Frame</i>	Bias solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>WeightsExportFormat</code> [OCHW]	<i>all Frame</i>	Weights import/export format. Can be <code>oCHW</code> or <code>OCHW</code> , with <code>o</code> the output feature map, <code>c</code> the input feature map (channel), <code>H</code> the kernel row and <code>w</code> the kernel column, in the order of the outermost dimension (in the leftmost position) to the innermost dimension (in the rightmost position)
<code>WeightsExportFlip</code> [0]	<i>all Frame</i>	If true, import/export flipped kernels

4.7.8 Pool

Pooling layer.

There are two CUDA models for this cell:

- `Frame_CUDA`, which uses CuDNN as back-end and only supports one-to-one input to output map connection;

- `Frame_EXT_CUDA`, which uses custom CUDA kernels and allows arbitrary connections between input and output maps (and can therefore be used to implement Maxout or both Maxout and Pooling simultaneously).

Maxout example In the following INI section, one implements a Maxout between each consecutive pair of 8 input maps:

```
[maxout_layer]
Input=...
Type=Pool
Model=Frame_EXT_CUDA
PoolWidth=1
PoolHeight=1
NbOutputs=4
Pooling=Max
Mapping.SizeY=2
Mapping.StrideY=2
```

The layer connectivity is the following:

# input map	1	■			
	2	■			
	3		■		
	4		■		
	5			■	
	6			■	
	7				■
	8				■
		1	2	3	4
		# output map			

Option [default value]	Description
Pooling	Type of pooling (Max or Average)
PoolWidth	Width of the pooling area
PoolHeight	Height of the pooling area
PoolDepth []	Depth of the pooling area (implies 3D pooling area)
OR	
PoolSize []	Pooling area size (implies 2D square pooling area)
OR	
PoolDims []	List of space-separated dimensions for N-D pooling area
NbOutputs	Number of output channels
StrideX [1]	X-axis stride of the pooling area
StrideY [1]	Y-axis stride of the pooling area
StrideZ []	Z-axis stride of the pooling area
OR	
Stride [1]	Stride of the pooling area
OR	
StrideDims []	List of space-separated stride dimensions for N-D pooling area
PaddingX [0]	X-axis input padding
PaddingY [0]	Y-axis input padding
PaddingZ []	Z-axis input padding
OR	
Padding [0]	Input padding

OR	
PaddingDims []	List of space-separated padding dimensions for N-D pooling area
ActivationFunction [Linear]	Activation function. Can be any of Logistic, LogisticWithLoss, Rectifier, Softplus, TanhLeCun, Linear, Saturation OR Tanh
Mapping.NbGroups []	Mapping: number of groups (mutually exclusive with all other Mapping.* options)
Mapping.ChannelsPerGroup []	Mapping: number of channels per group (mutually exclusive with all other Mapping.* options)
Mapping.SizeX [1]	Mapping canvas pattern default width
Mapping.SizeY [1]	Mapping canvas pattern default height
Mapping.Size [1]	Mapping canvas pattern default size (mutually exclusive with Mapping.SizeX and Mapping.SizeY)
Mapping.StrideX [1]	Mapping canvas default X-axis step
Mapping.StrideY [1]	Mapping canvas default Y-axis step
Mapping.Stride [1]	Mapping canvas default step (mutually exclusive with Mapping.StrideX and Mapping.StrideY)
Mapping.OffsetX [0]	Mapping canvas default X-axis offset
Mapping.OffsetY [0]	Mapping canvas default Y-axis offset
Mapping.Offset [0]	Mapping canvas default offset (mutually exclusive with Mapping.OffsetX and Mapping.OffsetY)
Mapping.NbIterations [0]	Mapping canvas pattern default number of iterations (0 means no limit)
Mapping(in).SizeX [1]	Mapping canvas pattern default width for input layer in
Mapping(in).SizeY [1]	Mapping canvas pattern default height for input layer in
Mapping(in).Size [1]	Mapping canvas pattern default size for input layer in (mutually exclusive with Mapping(in).SizeX and Mapping(in).SizeY)
Mapping(in).StrideX [1]	Mapping canvas default X-axis step for input layer in
Mapping(in).StrideY [1]	Mapping canvas default Y-axis step for input layer in
Mapping(in).Stride [1]	Mapping canvas default step for input layer in (mutually exclusive with Mapping(in).StrideX and Mapping(in).StrideY)
Mapping(in).OffsetX [0]	Mapping canvas default X-axis offset for input layer in
Mapping(in).OffsetY [0]	Mapping canvas default Y-axis offset for input layer in
Mapping(in).Offset [0]	Mapping canvas default offset for input layer in (mutually exclusive with Mapping(in).OffsetX and Mapping(in).OffsetY)
Mapping(in).NbIterations [0]	Mapping canvas pattern default number of iterations for input layer in (0 means no limit)

Configuration parameters (*Spike* models)

Option [default value]	Model(s)	Description
IncomingDelay [1 TimePs ;100 TimeFs] value	<i>all Spike</i>	Synaptic incoming delay w_{delay}

4.7.9 Unpool

Unpooling layer.

Option [default value]	Description
Pooling	Type of pooling (Max or Average)
PoolWidth	Width of the pooling area
PoolHeight	Height of the pooling area
PoolDepth []	Depth of the pooling area (implies 3D pooling area)
OR	
PoolSize []	Pooling area size (implies 2D square pooling area)
OR	
PoolDims []	List of space-separated dimensions for N-D pooling area
NbOutputs	Number of output channels
ArgMax	Name of the associated pool layer for the argmax (the pool layer input and the unpool layer output dimension must match)
StrideX [1]	X-axis stride of the pooling area
StrideY [1]	Y-axis stride of the pooling area
StrideZ []	Z-axis stride of the pooling area
OR	
Stride [1]	Stride of the pooling area
OR	
StrideDims []	List of space-separated stride dimensions for N-D pooling area
PaddingX [0]	X-axis input padding
PaddingY [0]	Y-axis input padding
PaddingZ []	Z-axis input padding
OR	
Padding [0]	Input padding
OR	
PaddingDims []	List of space-separated padding dimensions for N-D pooling area
ActivationFunction [Linear]	Activation function. Can be any of Logistic , LogisticWithLoss , Rectifier , Softplus , TanhLeCun , Linear , Saturation OR Tanh
Mapping.NbGroups []	Mapping: number of groups (mutually exclusive with all other Mapping.* options)
Mapping.ChannelsPerGroup []	Mapping: number of channels per group (mutually exclusive with all other Mapping.* options)
Mapping.SizeX [1]	Mapping canvas pattern default width
Mapping.SizeY [1]	Mapping canvas pattern default height
Mapping.Size [1]	Mapping canvas pattern default size (mutually exclusive with Mapping.SizeX and Mapping.SizeY)
Mapping.StrideX [1]	Mapping canvas default X-axis step
Mapping.StrideY [1]	Mapping canvas default Y-axis step
Mapping.Stride [1]	Mapping canvas default step (mutually exclusive with Mapping.StrideX and Mapping.StrideY)
Mapping.OffsetX [0]	Mapping canvas default X-axis offset
Mapping.OffsetY [0]	Mapping canvas default Y-axis offset
Mapping.Offset [0]	Mapping canvas default offset (mutually exclusive with Mapping.OffsetX and Mapping.OffsetY)

Mapping.NbIterations [0]	Mapping canvas pattern default number of iterations (0 means no limit)
Mapping(in).SizeX [1]	Mapping canvas pattern default width for input layer in
Mapping(in).SizeY [1]	Mapping canvas pattern default height for input layer in
Mapping(in).Size [1]	Mapping canvas pattern default size for input layer in (mutually exclusive with Mapping(in).SizeX and Mapping(in).SizeY)
Mapping(in).StrideX [1]	Mapping canvas default X-axis step for input layer in
Mapping(in).StrideY [1]	Mapping canvas default Y-axis step for input layer in
Mapping(in).Stride [1]	Mapping canvas default step for input layer in (mutually exclusive with Mapping(in).StrideX and Mapping(in).StrideY)
Mapping(in).OffsetX [0]	Mapping canvas default X-axis offset for input layer in
Mapping(in).OffsetY [0]	Mapping canvas default Y-axis offset for input layer in
Mapping(in).Offset [0]	Mapping canvas default offset for input layer in (mutually exclusive with Mapping(in).OffsetX and Mapping(in).OffsetY)
Mapping(in).NbIterations [0]	Mapping canvas pattern default number of iterations for input layer in (0 means no limit)

4.7.10 ElemWise

Element-wise operation layer.

Option [default value]	Description
NbOutputs	Number of output neurons
Operation	Type of operation (Sum, AbsSum, EuclideanSum, Prod, or Max)
Weights [1.0]	Weights for the Sum, AbsSum, and EuclideanSum operation, in the same order as the inputs
Shifts [0.0]	Shifts for the Sum and EuclideanSum operation, in the same order as the inputs
ActivationFunction [Linear]	Activation function. Can be any of Logistic, LogisticWithLoss, Rectifier, Softplus, TanhLeCun, Linear, Saturation Or Tanh

Given N input tensors T_i , performs the following operation:

Sum operation $T_{out} = \sum_1^N (w_i T_i + s_i)$

AbsSum operation $T_{out} = \sum_1^N (w_i |T_i|)$

EuclideanSum operation $T_{out} = \sqrt{\sum_1^N (w_i T_i + s_i)^2}$

Prod operation $T_{out} = \prod_1^N (T_i)$

Max operation $T_{out} = MAX_1^N (T_i)$

Examples Sum of two inputs ($T_{out} = T_1 + T_2$):

```
[elemwise_sum]
Input=layer1,layer2
Type=ElemWise
NbOutputs=[layer1]NbOutputs
Operation=Sum
```

Weighted sum of two inputs, by a factor 0.5 for `layer1` and 1.0 for `layer2` ($T_{out} = 0.5 \times T_1 + 1.0 \times T_2$):

```
[elemwise_weighted_sum]
Input=layer1,layer2
Type=ElemWise
NbOutputs=[layer1]NbOutputs
Operation=Sum
Weights=0.5 1.0
```

Single input scaling by a factor 0.5 and shifted by 0.1 ($T_{out} = 0.5 \times T_1 + 0.1$):

```
[elemwise_scale]
Input=layer1
Type=ElemWise
NbOutputs=[layer1]NbOutputs
Operation=Sum
Weights=0.5
Shifts=0.1
```

Absolute value of an input ($T_{out} = |T_1|$):

```
[elemwise_abs]
Input=layer1
Type=ElemWise
NbOutputs=[layer1]NbOutputs
Operation=Abs
```

4.7.11 FMP

Fractional max pooling layer ([Graham, 2014](#)).

Option [default value]	Description
NbOutputs	Number of output channels
ScalingRatio	Scaling ratio. The output size is $round\left(\frac{\text{input size}}{\text{scaling ratio}}\right)$.
ActivationFunction [Linear]	Activation function. Can be any of <code>Logistic</code> , <code>LogisticWithLoss</code> , <code>Rectifier</code> , <code>Softplus</code> , <code>TanhLeCun</code> , <code>Linear</code> , <code>Saturation</code> or <code>Tanh</code>

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
Overlapping [1]	<i>all Frame</i>	If true, use overlapping regions, else use disjoint regions
PseudoRandom [1]	<i>all Frame</i>	If true, use pseudorandom sequences, else use random sequences

4.7.12 Fc

Fully connected layer.

Option [default value]	Description
NbOutputs	Number of output neurons
WeightsFiller [NormalFiller(0.0, 0.05)]	Weights initial values filler
BiasFiller [NormalFiller(0.0, 0.05)]	Biases initial values filler
ActivationFunction [Tanh]	Activation function. Can be any of Logistic, LogisticWithLoss, Rectifier, Softplus, TanhLeCun, Linear, Saturation Or Tanh

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
NoBias [0]	<i>all Frame</i>	If true, don't use bias
BackPropagate [1]	<i>all Frame</i>	If true, enable backpropagation
Solvers.*	<i>all Frame</i>	Any solver parameters
WeightsSolver.*	<i>all Frame</i>	Weights solver parameters, take precedence over the Solvers.* parameters
BiasSolver.*	<i>all Frame</i>	Bias solver parameters, take precedence over the Solvers.* parameters
DropConnect [1.0]	Frame	If below 1.0, fraction of synapses that are disabled with drop connect

Configuration parameters (*Spike* models)

Option [default value]	Model(s)	Description
IncomingDelay [1 TimePs ;100 TimeFs]	<i>all Spike</i>	Synaptic incoming delay w_{delay}
Threshold [1.0]	Spike, Spike_RRAM	Threshold of the neuron I_{thres}
BipolarThreshold [1]	Spike, Spike_RRAM	If true, the threshold is also applied to the absolute value of negative values (generating negative spikes)
Leak [0.0]	Spike, Spike_RRAM	Neural leak time constant τ_{leak} (if 0, no leak)
Refractory [0.0]	Spike, Spike_RRAM	Neural refractory period T_{refrac}
TerminateDelta [0]	Spike, Spike_RRAM	Terminate delta
WeightsRelInit [0.0;0.05]	Spike	Relative initial synaptic weight w_{init}
WeightsMinMean [1;0.1]	Spike_RRAM	Mean minimum synaptic weight w_{min}
WeightsMaxMean [100;10.0]	Spike_RRAM	Mean maximum synaptic weight w_{max}
WeightsMinVarSlope [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMinVarOrigin [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMaxVarSlope [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsMaxVarOrigin [0.0]	Spike_RRAM	OxRAM specific parameter
WeightsSetProba [1.0]	Spike_RRAM	Intrinsic SET switching probability P_{SET} (upon receiving a SET programming pulse). Assuming uniform statistical distribution (not well supported by experiments on RRAM)

WeightsResetProba [1.0]	Spike_RRAM	Intrinsic RESET switching probability P_{RESET} (upon receiving a RESET programming pulse). Assuming uniform statistical distribution (not well supported by experiments on RRAM)
SynapticRedundancy [1]	Spike_RRAM	Synaptic redundancy (number of RRAM device per synapse)
BipolarWeights [0]	Spike_RRAM	Bipolar weights
BipolarIntegration [0]	Spike_RRAM	Bipolar integration
LtpProba [0.2]	Spike_RRAM	Extrinsic STDP LTP probability (cumulative with intrinsic SET switching probability P_{SET})
LtdProba [0.1]	Spike_RRAM	Extrinsic STDP LTD probability (cumulative with intrinsic RESET switching probability P_{RESET})
Stdpltp [1000 TimePs]	Spike_RRAM	STDP LTP time window T_{LTP}
InhibitRefractory [0 TimePs]	Spike_RRAM	Neural lateral inhibition period $T_{inhibit}$
EnableStdpl [1]	Spike_RRAM	If false, STDP is disabled (no synaptic weight change)
RefractoryIntegration [1]	Spike_RRAM	If true, reset the integration to 0 during the refractory period
DigitalIntegration [0]	Spike_RRAM	If false, the analog value of the devices is integrated, instead of their binary value

N2D2 IP only 4.7.13 Rbf

Radial basis function fully connected layer.

Option [default value]	Description
NbOutputs	Number of output neurons
CentersFiller [NormalFiller(0.5, 0.05)]	Centers initial values filler
ScalingFiller [NormalFiller(10.0, 0.05)]	Scaling initial values filler

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
Solvers.*	<i>all Frame</i>	Any solver parameters
CentersSolver.*	<i>all Frame</i>	Centers solver parameters, take precedence over the Solvers.* parameters
ScalingSolver.*	<i>all Frame</i>	Scaling solver parameters, take precedence over the Solvers.* parameters
RbfApprox [None]	Frame	Approximation for the Gaussian function, can be any of: None, Rectangular or SemiLinear

4.7.14 Softmax

Softmax layer.

Option [default value]	Description
NbOutputs	Number of output neurons
WithLoss [0]	Softmax followed with a multinomial logistic layer
GroupSize [0]	Softmax is applied on groups of outputs. The group size must be a divisor of NbOutputs parameter.

The softmax function performs the following operation, with $a_{x,y}^i$ and $b_{x,y}^i$ the input and the output respectively at position (x, y) on channel i :

$$b_{x,y}^i = \frac{\exp(a_{x,y}^i)}{\sum_{j=0}^N \exp(a_{x,y}^j)}$$

and

$$da_{x,y}^i = \sum_{j=0}^N (\delta_{ij} - a_{x,y}^i) a_{x,y}^j db_{x,y}^j$$

When the **WithLoss** option is enabled, compute the gradient directly in respect of the cross-entropy loss:

$$L_{x,y} = \sum_{j=0}^N t_{x,y}^j \log(b_{x,y}^j)$$

In this case, the gradient output becomes:

$$da_{x,y}^i = db_{x,y}^i$$

with

$$db_{x,y}^i = t_{x,y}^i - b_{x,y}^i$$

4.7.15 LRN

Local Response Normalization (LRN) layer.

Option [default value]	Description
NbOutputs	Number of output neurons

The response-normalized activity $b_{x,y}^i$ is given by the expression:

$$b_{x,y}^i = \frac{a_{x,y}^i}{\left(k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta}$$

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
N [5]	all <i>Frame</i>	Normalization window width in elements
Alpha [1.0e-4]	all <i>Frame</i>	Value of the alpha variance scaling parameter in the normalization formula
Beta [0.75]	all <i>Frame</i>	Value of the beta power parameter in the normalization formula
K [2.0]	all <i>Frame</i>	Value of the k parameter in normalization formula

4.7.16 LSTM

Long Short Term Memory Layer ([Hochreiter and Schmidhuber, 1997](#)).

Global layer parameters (*Frame_CUDA* models)

Option [default value]	Description
SeqLength	Maximum sequence length that the LSTM can take as an input.
BatchSize	Number of sequences used for a single weights actualisation process : size of the batch.
InputDim	Dimension of every element composing a sequence.
HiddenSize	Dimension of the LSTM inner state and output.
SingleBackpropFeeding [1]	If disabled return the full output sequence.
Bidirectional [0]	If enabled, build a bidirectional structure.
AllGatesWeightsFiller	All Gates weights initial values filler.
AllGatesBiasFiller	All Gates bias initial values filler.
WeightsInputGateFiller	Input gate previous layer and recurrent weights initial values filler. Take precedence over AllGatesWeightsFiller parameter.
WeightsForgetGateFiller	Forget gate previous layer and recurrent weights initial values filler. Take precedence over AllGatesWeightsFiller parameter.
WeightsCellGateFiller	Cell gate (or new memory) previous layer and recurrent weights initial values filler. Take precedence over AllGatesWeightsFiller parameter.
WeightsOutputGateFiller	Output gate previous layer and recurrent weights initial values filler. Take precedence over AllGatesWeightsFiller parameter.
BiasInputGateFiller	Input gate previous layer and recurrent bias initial values filler. Take precedence over AllGatesBiasFiller parameter.
BiasRecurrentForgetGateFiller	Forget gate recurrent bias initial values filler. Take precedence over AllGatesBiasFiller parameter. Often set to 1.0 to show better convergence performance.
BiasPreviousLayerForgetGateFiller	Forget gate previous layer bias initial values filler. Take precedence over AllGatesBiasFiller parameter.
BiasCellGateFiller	Cell gate (or new memory) previous layer and recurrent bias initial values filler. Take precedence over AllGatesBiasFiller parameter.
BiasOutputGateFiller	Output gate previous layer and recurrent bias initial values filler. Take precedence over AllGatesBiasFiller parameter.

HxFiller	Recurrent previous state initialisation. Often set to 0.0
CxFiller	Recurrent previous LSTM inner state initialisation. Often set to 0.0

Configuration parameters (*Frame_CUDA* models)

Option [default value]	Model(s)	Description
Solvers.*	<i>all Frame</i>	Any solver parameters
Dropout [0.0]	<i>all Frame</i>	The probability with which the value from input would be dropped.
InputMode []	<i>all Frame</i>	If enabled, drop the matrix multiplication of the input data.
Algo [0]	<i>all Frame</i>	Allow to choose different cuDNN implementation. Can be 0 : STANDARD, 1 : STATIC, 2 : DYNAMIC. Case 1 and 2 aren't supported yet.

Current restrictions :

- Only *Frame_Cuda* version is supported yet.
- The implementation only support input sequences with a fixed length associated with a single label.
- CuDNN structures requires the input data to be ordered as [1, InputDim, BatchSize, SeqLength]. Depending on the use case (like sequential-MNIST), the input data would need to be shuffled between the stimuli provisder and the RNN in order to process batches of data. No shuffling layer is yet operational. In that case, set batch to one for first experiments.

Further development requirements :

When it comes to RNN, two main factors needs to be considered to build proper interfaces :

1. Whether the input data has a variable or a fixed length over the data base, that is to say whether the input data will have a variable or fixed Sequence length. Of course the main strength of a RNN is to process variable length data.
2. Labelling granularity of the input data, that is to say wheteher every elements of a sequence is labelled or the sequence itself has only one label.

For instance, let's consider sentences as sequences of words in which every word would be part of a vocabulary. Sentences could have a variable length and every element/word would have a label. In that case, every relevant element of the output sequence from the recurrent structure is turned into a prediction throught a fully connected layer with a linear activation fonction and a softmax.

On the opposite, using sequential-MNIST database, the sequence length would be the same regarding every image and there is only one label for an image. In that case, the last element of the output sequence is the most relevant one to be turned into a prediction as it carries the information of the entire input sequence.

To provide flexibility according to these factors, the first implementation choice is to set a maximum sequence length `emphSeqLength` as an hyperparameter that the User provide. Variable

length sequences can be processed by padding the remaining steps of the input sequence. Then two cases occur as the labeling granularity is scaled at each element of the sequence or scaled at the sequence itself:

1. The sequence itself has only one label : The model has a fixed size with one fully connected




Figure 15: RNN model : variable sequence length and labeling scaled at the sequence

mapped to the relevant element of the output sequence according to the input sequence.

2. Every elements of a sequence is labelled :




Figure 16: RNN model : variable sequence length and labeling scaled at each element of the sequence

The model has a fixed size with one big fully connected (or T_{max} fully connected) mapped to the relevant elements of the output sequence according to the input sequence. The remaining elements need to be masked so it doesn't influence longer sequences.

Development guidance :

- Replace the inner local variables of LSTMCell_Frame_Cuda with a generic layer of shuffling (on device) to enable the the process of data batch.

- Develop some kind of label embedding within the layer to better articulate the labeling granularity of the input data.
- Adapt structures to support the STATIC and DYNAMIC algorithm of cuDNN functions.

4.7.17 Dropout

Dropout layer ([Srivastava et al., 2012](#)).

Option [default value]	Description
NbOutputs	Number of output neurons

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
Dropout [0.5]	all Frame	The probability with which the value from input would be dropped

4.7.18 Padding

Padding layer.

Option [default value]	Description
NbOutputs	Number of output neurons
TopPadding	Size of the top padding (positive or negative)
BottomPadding	Size of the bottom padding (positive or negative)
LeftPadding	Size of the left padding (positive or negative)
RightPadding	Size of the right padding (positive or negative)

The padding layer allow to insert asymmetric padding for each layer axes.

4.7.19 Resize

Resize layer can be applied to change dimension of features maps or of stimuli provider.

Option [default value]	Description
NbOutputs	Number of output feature maps
OutputHeight	Output height dimension
OutputWidth	Output width dimension
Mode	Resize interpolation mode. Can be, Bilinear or BilinearTF (TensorFlow implementation)

Configuration parameters

Option [default value]	Model(s)	Description
<code>AlignCorners</code> [True]	<i>all Frame</i>	Corner alignment mode if <code>BilinearTF</code> is used as interpolation mode

4.7.20 BatchNorm

Batch Normalization layer ([Ioffe and Szegedy, 2015](#)).

Option [default value]	Description
NbOutputs	Number of output neurons
<code>ActivationFunction</code> [Tanh]	Activation function. Can be any of <code>Logistic</code> , <code>LogisticWithLoss</code> , <code>Rectifier</code> , <code>Softplus</code> , <code>TanhLeCun</code> , <code>Linear</code> , <code>Saturation</code> Or <code>Tanh</code>
<code>ScalesSharing</code> []	Share the scales with an other layer
<code>BiasesSharing</code> []	Share the biases with an other layer
<code>MeansSharing</code> []	Share the means with an other layer
<code>VariancesSharing</code> []	Share the variances with an other layer

Configuration parameters (*Frame* models)

Option [default value]	Model(s)	Description
<code>Solvers.*</code>	<i>all Frame</i>	Any solver parameters
<code>ScaleSolver.*</code>	<i>all Frame</i>	Scale solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>BiasSolver.*</code>	<i>all Frame</i>	Bias solver parameters, take precedence over the <code>Solvers.*</code> parameters
<code>Epsilon</code> [0.0]	<i>all Frame</i>	Epsilon value used in the batch normalization formula. If 0.0, automatically choose the minimum possible value.
<code>MovingAverageMomentum</code> [0.1]	<i>all Frame</i>	<code>MovingAverageMomentum</code> : used for the moving average of batch-wise means and standard deviations during training. The closer to 1.0, the more it will depend on the last batch.

4.7.21 Transformation

Transformation layer, which can apply any transformation described in 4.6.1. Useful for fully CNN post-processing for example.

Option [default value]	Description
NbOutputs	Number of outputs
Transformation	Name of the transformation to apply

The Transformation options must be placed in the same section.
Usage example for fully CNNs:

```
[post.Transformation-thres]
Input=... ; for example, network's logistic of softmax output layer
NbOutputs=1
Type=Transformation
Transformation=ThresholdTransformation
Operation=ToZero
Threshold=0.75

[post.Transformation-morpho]
Input=post.Transformation-thres
NbOutputs=1
Type=Transformation
Transformation=MorphologyTransformation
Operation=Opening
Size=3
```

5 Tutorials

5.1 Learning deep neural networks: tips and tricks

5.1.1 Choose the learning solver

Generally, you should use the SGD solver with a momentum (typical value for the momentum: 0.9). It generalizes better, often significantly better, than adaptive methods like Adam ([Wilson et al., 2017](#)).

Adaptive solvers, like Adam, may be used for fast exploration and prototyping, thanks to their fast convergence.

5.1.2 Choose the learning hyper-parameters

You can use the `-find-lr` option available in the `n2d2` executable to automatically find the best learning rate for a given neural network.

Usage example:

```
./n2d2 model.ini -find-lr 10000
```

This command starts from a very low learning rate ($1.0e-6$) and increase it exponentially to reach the maximum value (10.0) after 10000 steps, as shown in figure 17. The loss change during this phase is then plotted in function of the learning rate, as shown in figure 18.




Figure 17: Exponential increase of the learning rate over the specified number of iterations, equals to the number of steps divided by the batch size (here: 24).




Figure 18: Loss change as a function of the learning rate.

Note that in N2D2, the learning rate is automatically normalized by the global batch size ($N \times \text{IterationSize}$) for the `SGDSolver`. A simple linear scaling rule is used, as recommended in (Goyal et al., 2017). The effective learning rate α_{eff} applied for parameters update is therefore:

$$\alpha_{\text{eff}} = \frac{\alpha}{N \times \text{IterationSize}} \text{ with } \alpha = \text{LearningRate}$$

Typical values for the `SGDSolver` are:

```
Solvers.LearningRate=0.01
Solvers.Decay=0.0001
Solvers.Momentum=0.9
```

5.1.3 Convergence and normalization

Deep networks (> 30 layers) and especially residual networks usually don't converge without normalization. Indeed, batch normalization is almost always used. *ZeroInit* is a method that can be used to overcome this issue without normalization (Zhang et al., 2019).

5.2 Building a classifier neural network

For this tutorial, we will use the classical MNIST handwritten digit dataset. A driver module already exists for this dataset, named `MNIST_IDX_Database`.

To instantiate it, just add the following lines in a new INI file:

```
[database]
Type=MNIST_IDX_Database
Validation=0.2 ; Use 20% of the dataset for validation
```

In order to create a neural network, we first need to define its input, which is declared with a `[sp]` section (*sp* for *StimuliProvider*). In this section, we configure the size of the input and the batch size:

```
[sp]
SizeX=32
SizeY=32
BatchSize=128
```

We can also add pre-processing transformations to the *StimuliProvider*, knowing that the final data size after transformations must match the size declared in the [sp] section. Here, we must rescale the MNIST 28x28 images to match the 32x32 network input size.

```
[sp.Transformation_1]
Type=RescaleTransformation
Width=[sp]SizeX
Height=[sp]SizeY
```

Next, we declare the neural network layers. In this example, we reproduced the well-known LeNet network. The first layer is a 5x5 convolutional layer, with 6 channels. Since there is only one input channel, there will be only 6 convolution kernels in this layer.

```
[conv1]
Input=sp
Type=Conv
KernelWidth=5
KernelHeight=5
NbOutputs=6
```

The next layer is a 2x2 MAX pooling layer, with a stride of 2 (non-overlapping MAX pooling).

```
[pool1]
Input=conv1
Type=Pool
PoolWidth=2
PoolHeight=2
NbOutputs=[conv1]NbOutputs
Stride=2
Pooling=Max
Mapping.Size=1 ; One to one connection between input and output channels
```

The next layer is a 5x5 convolutional layer with 16 channels.

```
[conv2]
Input=pool1
Type=Conv
KernelWidth=5
KernelHeight=5
NbOutputs=16
```

Note that in LeNet, the [conv2] layer is not fully connected to the pooling layer. In N2D2, a custom mapping can be defined for each input connection. The connection of n -th output map to the inputs is defined by the n -th column of the matrix below, where the rows correspond to the inputs.

```
Mapping(pool1)=\
1 0 0 0 1 1 1 0 0 1 1 1 1 0 1 1 \
1 1 0 0 0 1 1 1 0 0 1 1 1 1 0 1 \
1 1 1 0 0 0 1 1 1 0 0 1 0 1 1 1 \
0 1 1 1 0 0 1 1 1 1 0 0 1 0 1 1 \
0 0 1 1 1 0 0 1 1 1 1 0 1 1 0 1 \
0 0 0 1 1 1 0 0 1 1 1 1 0 1 1 1
```

Another MAX pooling and convolution layer follow:

```
[pool2]
Input=conv2
Type=Pool
PoolWidth=2
PoolHeight=2
```

```
NbOutputs=[conv2]NbOutputs
Stride=2
Pooling=Max
Mapping.Size=1
```

```
[conv3]
Input=pool2
Type=Conv
KernelWidth=5
KernelHeight=5
NbOutputs=120
```

The network is composed of two fully-connected layers of 84 and 10 neurons respectively:

```
[fc1]
Input=conv3
Type=Fc
NbOutputs=84
```

```
[fc2]
Input=fc1
Type=Fc
NbOutputs=10
```

Finally, we use a softmax layer to obtain output classification probabilities and compute the loss function.

```
[softmax]
Input=fc2
Type=Softmax
NbOutputs=[fc2]NbOutputs
WithLoss=1
```

In order to tell N2D2 to compute the error and the classification score on this softmax layer, one must attach a N2D2 *Target* to this layer, with a section with the same name suffixed with *.Target*:

```
[softmax.Target]
```

By default, the activation function for the convolution and the fully-connected layers is the hyperbolic tangent. Because the [fc2] layer is fed to a softmax, it should not have any activation function. We can specify it by adding the following line in the [fc2] section:

```
[fc2]
...
ActivationFunction=Linear
```

In order to improve further the networks performances, several things can be done:

- **Use ReLU activation functions.** In order to do so, just add the following in the [conv1], [conv2], [conv3] and [fc1] layer sections:

```
ActivationFunction=Rectifier
```

For the ReLU activation function to be effective, the weights must be initialized carefully, in order to avoid dead units that would be stuck in the $]-\infty, 0]$ output range before the ReLU function. In N2D2, one can use a custom `WeightsFiller` for the weights initialization. For the ReLU activation function, a popular and efficient filler is the so-called `XavierFiller` (see the 4.7.2 section for more information):

```
WeightsFiller=XavierFiller
```

- **Use dropout layers.** Dropout is highly effective to improve the network generalization capacity. Here is an example of a dropout layer inserted between the [fc1] and [fc2] layers:

```
[fc1]
...
```

```
[fc1.drop]
Input=fc1
Type=Dropout
NbOutputs=[fc1]NbOutputs

[fc2]
Input=fc1.drop ; Replaces "Input=fc1"
...
```

- **Tune the learning parameters.** You may want to tune the learning rate and other learning parameters depending on the learning problem at hand. In order to do so, you can add a configuration section that can be common (or not) to all the layers. Here is an example of configuration section:

```
[conv1]
...
ConfigSection=common.config

[...]
...

[common.config]
NoBias=1
WeightsSolver.LearningRate=0.05
WeightsSolver.Decay=0.0005
Solvers.LearningRatePolicy=StepDecay
Solvers.LearningRateStepSize=[sp]_EpochSize
Solvers.LearningRateDecay=0.993
Solvers.Clamping=-1.0:1.0
```

For more details on the configuration parameters for the `Solver`, see section 4.7.3.

- **Add input distortion.** See for example the `DistortionTransformation` (section 4.6.1).

The complete INI model corresponding to this tutorial can be found in *models/LeNet.ini*.

In order to use CUDA/GPU accelerated learning, the default layer model should be switched to `Frame_CUDA`. You can enable this model by adding the following line at the top of the INI file (before the first section):

```
DefaultModel=Frame_CUDA
```

5.3 Building a segmentation neural network

In this tutorial, we will learn how to do image segmentation with N2D2. As an example, we will implement a face detection and gender recognition neural network, using the IMDB-WIKI dataset.

First, we need to instantiate the IMDB-WIKI dataset built-in N2D2 driver:

```
[database]
Type=IMDBWIKI_Database
WikiSet=1 ; Use the WIKI part of the dataset
IMDBSet=0 ; Don't use the IMDB part (less accurate annotation)
Learn=0.90
Validation=0.05
DefaultLabel=background ; Label for pixels outside any ROI (default is no label, pixels are ignored)
```

We must specify a default label for the background, because we want to learn to differentiate faces from the background (and not simply ignore the background for the learning).

The network input is then declared:

```
[sp]
SizeX=480
```

```
SizeY=360
BatchSize=48
CompositeStimuli=1
```

In order to work with segmented data, i.e. data with bounding box annotations or pixel-wise annotations (as opposed to a single label per data), one must enable the `CompositeStimuli` option in the `[sp]` section.

We can then perform various operations on the data before feeding it to the network, like for example converting the 3-channels RGB input images to single-channel gray images:

```
[sp.Transformation-1]
Type=ChannelExtractionTransformation
CSChannel=Gray
```

We must only rescale the images to match the networks input size. This can be done using a `RescaleTransformation`, followed by a `PadCropTransformation` if one want to keep the images aspect ratio.

```
[sp.Transformation-2]
Type=RescaleTransformation
Width=[sp]SizeX
Height=[sp]SizeY
KeepAspectRatio=1 ; Keep images aspect ratio

; Required to ensure all the images are the same size
[sp.Transformation-3]
Type=PadCropTransformation
Width=[sp]SizeX
Height=[sp]SizeY
```

A common additional operation to extend the learning set is to apply random horizontal mirror to images. This can be achieved with the following `FlipTransformation`:

```
[sp.OnTheFlyTransformation-4]
Type=FlipTransformation
RandomHorizontalFlip=1
ApplyTo=LearnOnly ; Apply this transformation only on the learning set
```

Note that this is an *on-the-fly* transformation, meaning it cannot be cached and is re-executed every time even for the same stimuli. We also apply this transformation only on the learning set, with the `ApplyTo` option.

Next, the neural network can be described:

```
[conv1.1]
Input=sp
Type=Conv
...

[pool1]
...

[...]
...

[fc2]
Input=drop1
Type=Conv
...

[drop2]
Input=fc2
Type=Dropout
NbOutputs=[fc2]NbOutputs
```

A full network description can be found in the *IMDBWIKI.ini* file in the *models* directory of N2D2. It is a fully-CNN network.

Here we will focus on the output layers required to detect the faces and classify their gender. We start from the [drop2] layer, which has 128 channels of size 60x45.

5.3.1 Faces detection

We want to first add an output stage for the faces detection. It is a 1x1 convolutional layer with a single 60x45 output map. For each output pixel, this layer outputs the probability that the pixel belongs to a face.

```
[fc3.face]
Input=drop2
Type=Conv
KernelWidth=1
KernelHeight=1
NbOutputs=1
Stride=1
ActivationFunction=LogisticWithLoss
WeightsFiller=XavierFiller
ConfigSection=common.config ; Same solver options that the other layers
```

In order to do so, the activation function of this layer must be of type `LogisticWithLoss`.

We must also tell N2D2 to compute the error and the classification score on this softmax layer, by attaching a N2D2 *Target* to this layer, with a section with the same name suffixed with *.Target*:

```
[fc3.face.Target]
LabelsMapping=${N2D2_MODELS}/IMDBWIKI_target_face.dat
; Visualization parameters
NoDisplayLabel=0
LabelsHueOffset=90
```

In this *Target*, we must specify how the dataset annotations are mapped to the layer's output. This can be done in a separate file using the `LabelsMapping` parameter. Here, since the output layer has a single output per pixel, the target value can only be 0 or 1. A target value of -1 means that this output is ignored (no error back-propagated). Since the only annotations in the IMDB-WIKI dataset are faces, the mapping described in the *IMDBWIKI_target_face.dat* file is easy:

```
# background
background 0

# padding (*) is ignored (-1)
* -1

# not background = face
default 1
```

5.3.2 Gender recognition

We can also add a second output stage for gender recognition. Like before, it would be a 1x1 convolutional layer with a single 60x45 output map. But here, for each output pixel, this layer would output the probability that the pixel represents a female face.

```
[fc3.gender]
Input=drop2
Type=Conv
KernelWidth=1
KernelHeight=1
NbOutputs=1
Stride=1
ActivationFunction=LogisticWithLoss
WeightsFiller=XavierFiller
```

```
ConfigSection=common.config
```

The output layer is therefore identical to the face's output layer, but the target mapping is different. For the target mapping, the idea is simply to ignore all pixels not belonging to a face and affect the target 0 to male pixels and the target 1 to female pixels.

```
[fc3.gender.Target]
LabelsMapping=${N2D2_MODELS}/IMDBWIKI_target_gender.dat
; Only display gender probability for pixels detected as face pixels
MaskLabelTarget=fc3.face.Target
MaskedLabel=1
```

The content of the *IMDBWIKI_target_gender.dat* file would therefore look like:

```
# background
# ?-* (unknown gender)
# padding
default -1

# male gender
M-? 0 # unknown age
M-0 0
M-1 0
M-2 0
...
M-98 0
M-99 0

# female gender
F-? 1 # unknown age
F-0 1
F-1 1
F-2 1
...
F-98 1
F-99 1
```

5.3.3 ROIs extraction

The next step would be to extract detected face ROIs and assign for each ROI the most probable gender. To this end, we can first set a detection threshold, in terms of probability, to select face pixels. In the following, the threshold is fixed to 75% face probability:

```
[post.Transformation-thres]
Input=fc3.face
Type=Transformation
NbOutputs=1
Transformation=ThresholdTransformation
Operation=ToZero
Threshold=0.75
```

We can then assign a target of type `TargetROIs` to this layer that will automatically create the bounding box using a segmentation algorithm.

```
[post.Transformation-thres.Target-face]
Type=TargetROIs
MinOverlap=0.33 ; Min. overlap fraction to match the ROI to an annotation
FilterMinWidth=5 ; Min. ROI width
FilterMinHeight=5 ; Min. ROI height
FilterMinAspectRatio=0.5 ; Min. ROI aspect ratio
FilterMaxAspectRatio=1.5 ; Max. ROI aspect ratio
LabelsMapping=${N2D2_MODELS}/IMDBWIKI_target_face.dat
```

In order to assign a gender to the extracted ROIs, the above target must be modified to:


```
[post.Transformation-thres.Target-gender]
Type=TargetROIs
ROIsLabelTarget=fc3.gender.Target
MinOverlap=0.33
FilterMinWidth=5
FilterMinHeight=5
FilterMinAspectRatio=0.5
FilterMaxAspectRatio=1.5
LabelsMapping=${N2D2_MODELS}/IMDBWIKI_target_gender.dat
```

Here, we use the `fc3.gender.Target` target to determine the most probable gender of the ROI.

5.3.4 Data visualization

For each *Target* in the network, a corresponding folder is created in the simulation directory, which contains learning, validation and test confusion matrixes. The output estimation of the network for each stimulus is also generated automatically for the test dataset and can be visualized with the `./test.py` helper tool. An example is shown in figure 19.

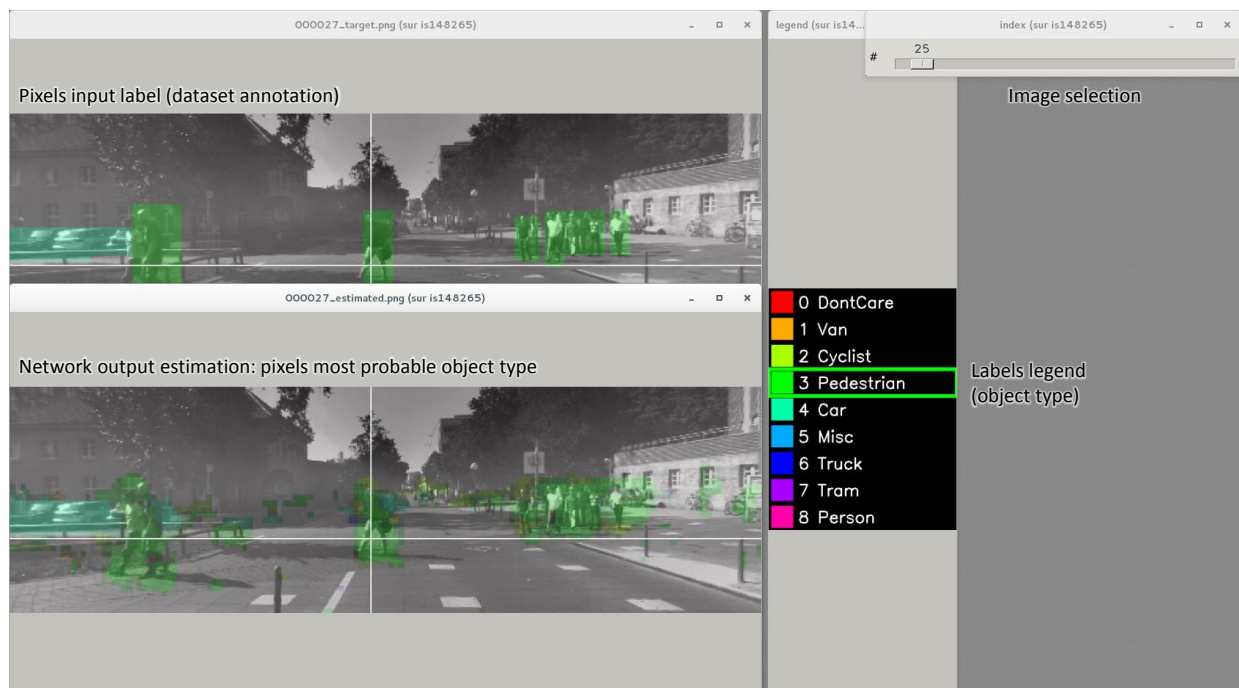


Figure 19: Example of the target visualization helper tool.

5.4 Transcoding a learned network in spike-coding

N2D2 embeds an event-based simulator (historically known as 'Xnet') and allows to transcode a whole DNN in a spike-coding version and evaluate the resulting spiking neural network performances. In this tutorial, we will transcode the LeNet network described in section 5.2.

5.4.1 Render the network compatible with spike simulations

The first step is to specify that we want to use a transcode model (allowing both formal and spike simulation of the same network), by changing the `DefaultModel` to:

```
DefaultModel=Transcode_CUDA
```

In order to perform spike simulations, the input of the network must be of type *Environment*, which is a derived class of *StimuliProvider* that adds spike coding support. In the INI model file, it

is therefore necessary to replace the `[sp]` section by an `[env]` section and replace all references of `sp` to `env`.

Note that these changes have at this point no impact at all on the formal coding simulations. The beginning of the INI file should be:

```
DefaultModel=Transcode_CUDA

; Database
[database]
Type=MNIST_IDX_Database
Validation=0.2 ; Use 20% of the dataset for validation

; Environment
[env]
SizeX=32
SizeY=32
BatchSize=128

[env.Transformation_1]
Type=RescaleTransformation
Width=[env]SizeX
Height=[env]SizeY

[conv1]
Input=env
...
```

The dropout layer has no equivalence in spike-coding inference and must be removed:

```
...
[fc1.dropout]
Input=fc1
Type=Dropout
NbOutputs=[fc1]NbOutputs

[fc2]
Input=fc1.dropout
...
```

The softmax layer has no equivalence in spike-coding inference and must be removed as well. The *Target* must therefore be attached to `[fc2]`:

```
...
[softmax]
Input=fc2
Type=Softmax
NbOutputs=[fc2]NbOutputs
WithLoss=1

[softmax.Target]

[fc2.Target]
...
```

The network is now compatible with spike-coding simulations. However, we did not specify at this point how to translate the input stimuli data into spikes, nor the spiking neuron parameters (threshold value, leak time constant...).

5.4.2 Configure spike-coding parameters

The first step is to configure how the input stimuli data must be coded into spikes. To this end, we must attach a configuration section to the *Environment*. Here, we specify a periodic coding with random initial jitter with a minimum period of 10 ns and a maximum period of 100 us:

```

[env]
...
ConfigSection=env.config

[env.config]
; Spike-based computing
StimulusType=JitteredPeriodic
PeriodMin=1,000,000 ; unit = fs
PeriodMeanMin=10,000,000 ; unit = fs
PeriodMeanMax=100,000,000,000 ; unit = fs
PeriodRelStdDev=0.0

```

The next step is to specify the neurons parameters, that will be common to all layers and can therefore be specified in the `[common.config]` section. In N2D2, the base spike-coding layers use a Leaky Integrate-and-Fire (LIF) neuron model. By default, the leak time constant is zero, resulting to simple Integrate-and-Fire (IF) neurons.

Here we simply specify that the neurons threshold must be the unity, that the threshold is only positive and that there is no incoming synaptic delay:

```

[common.config]
...
; Spike-based computing
Threshold=1.0
BipolarThreshold=0
IncomingDelay=0

```

Finally, we can limit the number of spikes required for the computation of each stimulus by adding a decision delta threshold at the output layer:

```

[fc2]
...
ConfigSection=common.config,fc2.config

[fc2.Target]

[fc2.config]
; Spike-based computing
TerminateDelta=4
BipolarThreshold=1

```

The complete INI model corresponding to this tutorial can be found in *models/LeNet_Spike.ini*.

Here is a summary of the steps required to reproduce the whole experiment:

```

./n2d2 "$N2D2_MODELS/LeNet.ini" -learn 6000000 -log 100000
./n2d2 "$N2D2_MODELS/LeNet_Spike.ini" -test

```

The final recognition rate reported at the end of the spike inference should be almost identical to the formal coding network (around 99% for the LeNet network).

Various statistics are available at the end of the spike-coding simulation in the *stats_spike* folder and the *stats_spike.log* file. Looking in the *stats_spike.log* file, one can read the following line towards the end of the file:

```
Read events per virtual synapse per pattern (average): 0.654124
```

This line reports the average number of accumulation operations per synapse per input stimulus in the network. If this number is below 1.0, it means that the spiking version of the network is more efficient than its formal counterpart in terms of total number of operations!

References

- M. Cordts, M. Omran, S. Ramos, T. Rehfeld, M. Enzweiler, R. Benenson, U. Franke, S. Roth, and B. Schiele. The cityscapes dataset for semantic urban scene understanding. In *Proc. of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- P. Dollár, C. Wojek, B. Schiele, and P. Perona. Pedestrian detection: A benchmark. In *CVPR*, 2009.
- L. Fei-Fei, R. Fergus, and P. Perona. Learning generative visual models from few training examples: an incremental bayesian approach tested on 101 object categories. In *IEEE. CVPR 2004, Workshop on Generative-Model Based Vision*, 2004.
- X. Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks. In *International conference on artificial intelligence and statistics*, page 249–256, 2010.
- P. Goyal, P. Dollár, R. B. Girshick, P. Noordhuis, L. Wesolowski, A. Kyrola, A. Tulloch, Y. Jia, and K. He. Accurate, large minibatch SGD: training imagenet in 1 hour. *CoRR*, abs/1706.02677, 2017. URL <http://arxiv.org/abs/1706.02677>.
- B. Graham. Fractional max-pooling. *CoRR*, abs/1412.6071, 2014.
- G. Griffin, A. Holub, and P. Perona. Caltech-256 object category dataset, 2007.
- K. He, X. Zhang, S. Ren, and J. Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *Proceedings of the 2015 IEEE International Conference on Computer Vision (ICCV)*, ICCV ’15, pages 1026–1034, 2015. doi: 10.1109/ICCV.2015.123.
- S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. doi: 10.1162/neco.1997.9.8.1735.
- S. Houben, J. Stallkamp, J. Salmen, M. Schlipsing, and C. Igel. Detection of traffic signs in real-world images: The German Traffic Sign Detection Benchmark. In *International Joint Conference on Neural Networks*, number 1288, 2013.
- S. Ioffe and C. Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *CoRR*, abs/1502.03167, 2015.
- V. Jain and E. Learned-Miller. FDDB: A benchmark for face detection in unconstrained settings, 2010.
- D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. *CoRR*, abs/1412.6980, 2014. URL <http://arxiv.org/abs/1412.6980>.
- A. Krizhevsky. Learning multiple layers of features from tiny images, 2009.
- Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. In *Proceedings of the IEEE*, volume 86, pages 2278–2324, 1998.
- P. Lucey, J. F. Cohn, T. Kanade, J. Saragih, Z. Ambadar, and I. Matthews. The Extended Cohn-Kanade Dataset (CK+): A complete dataset for action unit and emotion-specified expression. 2010.
- A. Rakotomamonjy and G. Gasso. Histogram of gradients of time-frequency representations for audio scene detection, 2014.

-
- O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpathy, A. Khosla, M. Bernstein, A. C. Berg, and L. Fei-Fei. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015. doi: 10.1007/s11263-015-0816-y.
- N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: A simple way to prevent neural networks from voverfitting. *Journal of Machine Learning Research*, 15: 1929–1958, 2012.
- J. Stallkamp, M. Schlipsing, J. Salmen, and C. Igel. Man vs. computer: Benchmarking machine learning algorithms for traffic sign recognition. *Neural Networks*, 2012. ISSN 0893-6080. doi: 10.1016/j.neunet.2012.02.016.
- J. W. Lockhart, G. Weiss, J. C. Xue, S. T. Gallagher, A. B. Grosner, and T. Pulickal. Design considerations for the wisdm smart phone-based sensor mining architecture. 01 2011. doi: 10.1145/2003653.2003656.
- P. Warden. Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition. *ArXiv e-prints*, Apr. 2018. URL <https://arxiv.org/abs/1804.03209>.
- A. C. Wilson, R. Roelofs, M. Stern, N. Srebro, and B. Recht. The Marginal Value of Adaptive Gradient Methods in Machine Learning. *arXiv e-prints*, art. arXiv:1705.08292, May 2017.
- G. Xia, X. Bai, J. Ding, Z. Zhu, S. J. Belongie, J. Luo, M. Datcu, M. Pelillo, and L. Zhang. DOTA: A large-scale dataset for object detection in aerial images. *CoRR*, abs/1711.10398, 2017. URL <http://arxiv.org/abs/1711.10398>.
- H. Zhang, Y. N. Dauphin, and T. Ma. Residual learning without normalization via better initialization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=H1gsz30cKX>.